

# What Survives the Translation from Brain to Von Neumann Machine: Learning Is Pruning, and Four Mechanisms That Did Not Transfer

Dongin Kang  
Independent Researcher  
shaloooooooooom@gmail.com

Draft, October 2026

## Abstract

A brain computes, stores and learns in the same synapses; a von Neumann machine separates processor from memory. Which brain mechanisms survive the translation? We ported five mechanisms and one follow-up to GPU training under one protocol: gradual synaptic pruning during learning, a fast hippocampus consolidating into a slow cortex during sleep, thalamic routing that decides what to process, predictive coding that skips the predictable, and spike-timing-dependent plasticity. One principle transfers. At matched final connection budgets, removing connections gradually while training beats training small, pruning once after training and dynamic sparse training (RigL) at tight budgets, ties with prune-after at loose ones, and its margin grows with sparsity: at 0.5% of the connections it reaches 97.8% on MNIST against 94.1%, 94.5% and 96.3%, and 89.0% on a CIFAR-10 ResNet-18 against 81.3%, 83.6% and 84.9%. The gain comes from the gradual removal, not from sparsity; it survives at matched inference FLOPs and costs 3–59× the training compute of the small network. A pre-trained start adds 3.5–5.3 points over pruning from scratch without saving adaptation compute (93.7% at 5% of an ImageNet ResNet-18’s connections) and becomes a liability at 0.5%. Local, activity-dependent pruning criteria trail magnitude pruning and collapse on CNNs. Four structures do not transfer: hippocampus–cortex sleep matches class-balanced replay; thalamic routing saves only 4–5% of total compute at 50–65 tokens; predictive coding neither picks which tokens to skip better than chance nor fills the gap better than leaving it; per-input subnetworks at inference are input-specific yet worse than a fixed mask. We offer compute–memory unity as the hypothesis behind this pattern, not a conclusion. Finally, an STDP-trained spiking network sparsifies itself without explicit pruning.

## 1 Introduction

A brain computes, stores and learns in the same synapses; a von Neumann machine keeps its processor and its memory apart and moves every weight between them. Most brain-inspired proposals for deep learning copy one brain mechanism into this second kind of machine and report whether it helps. We ask the question the other way round: when the same protocol and a common set of metrics are applied to each of several brain mechanisms, which of them survive the translation, and which do not?

Synaptic density in human cortex peaks in early childhood and falls by roughly half by adulthood [14]; the pruning is activity dependent and interleaved with learning. Deep learning takes the opposite route: parameters are accumulated from random initialisation, and sparsity, when wanted, is imposed afterwards [12, 13]. The lottery-ticket literature showed that large networks contain small trainable subnetworks [8], and gradual magnitude pruning [28] and dynamic sparse training

[7, 19] showed that sparsity can be reached during training. What has not been laid out side by side, at identical final budgets and with the full cost accounting, is the simple developmental question: does it matter whether a network is pruned *while* it learns, after it learns, or never grown large in the first place?

This paper is the first empirical instalment of a broader programme that set out to test four brain principles (pruning, compute–memory unity, retrieval structure and information concentration) in small models, by porting five concrete mechanisms to ordinary GPU training: gradual synaptic pruning during learning, a fast hippocampus consolidating into a slow cortex during sleep, thalamic routing that decides what to process, predictive coding that skips the predictable, and spike-timing-dependent plasticity, plus one follow-up suggested by the first result, per-input pruning at inference time. Of the four principles only pruning is tested directly: compute–memory unity cannot be tested on a GPU and appears only as a hypothesis in Section 6, and retrieval structure and information concentration motivated the routing and attention experiments. Table 5 maps each mechanism to its reduction and outcome. We report the one principle that survived contact with the data, and the rest honestly as negative results, because they constrain what a brain-inspired architecture should and should not borrow.

## Contributions.

1. **A matched-budget comparison of four routes to sparsity** (Section 4). On MNIST (MLP, five budgets from 10% to 0.5% of a 1.86M-connection network) and CIFAR-10 (VGG-style CNN, three budgets from 10% to 1% of 2.2M connections), prune-during with a global magnitude criterion is the most accurate route at every budget of 2% and below, ties with prune-after (within 0.4 points) at the loosest budgets, and its margin over the additive baseline grows from 0.3 to 3.7 points (MNIST) and from 3.0 to 8.6 points (CIFAR-10) as the budget shrinks. A ResNet-18 replication (11.2M weights, three budgets from 5% to 0.5%) reproduces the winner and the growth of the margin, with prune-during at 0.5% reaching 89.0% against 81.3% (dense small), 83.6% (prune-after) and 84.9% (RigL).
2. **The benefit is in the gradual removal, not in sparsity.** Pruning the same network once after training collapses accuracy at the pruning step (to below 60% at 0.5% density on MNIST) and recovers only partially; gradual removal with the global criterion never falls below 97% (Figure 1).
3. **Honest cost accounting.** Prune-during costs 3–59× the cumulative training FLOPs of the dense small network because it starts large; RigL costs about the same as the dense small network on the MLP and CNN and 2.3–2.5× on ResNet-18. On CNNs, equal connection counts are not equal FLOPs; constraining the layer-wise allocation to the ERK pattern, which uses fewer inference FLOPs than the dense small network, still leaves a 4.2-point advantage at 1%.
4. **Transfer from a pre-trained network** (Table 4). Pruning an ImageNet-pre-trained ResNet-18 *during* its adaptation to CIFAR-10 keeps 93.7% at 5% and 91.0% at 2% of its connections, against 88.4% and 87.5% from scratch and 86.9% for RigL on the same weights; at 0.5% the advantage reverses and both one-shot pruning and training from scratch win: the inherited start adds accuracy, not adaptation compute, and pays off only down to about 2%.
5. **Local pruning rules.** A Hebbian correlation rule beats random removal but falls below the dense small network at every budget under 10%. A synapse-level rule (“synaptic drive”: the

input flowing through a synapse when its post-synaptic unit fires) is competitive on the MLP but collapses on the CNN through channel death unless normalised per neuron, and trails magnitude pruning in both settings.

6. **Negative results** (Section 5) for a simple sleep-consolidation scheme with a fast hippocampal learner, for predictive token skipping (neither its choice of tokens nor its predicted-change substitution beats a random or identity control), for thalamic routing (pre-routed sparse attention), and for input-dependent pruning at inference time (which forms input-specific subnetworks yet loses to a fixed mask at the same budget), together with the observation that STDP alone sparsifies a spiking network during learning.

## 2 Related work

**Pruning.** Magnitude pruning after training [12], gradual pruning during training with a cubic sparsity schedule [28], the lottery-ticket hypothesis [8] and data-free saliency criteria such as SynFlow [24] are the closest antecedents; Hoefler et al. [13] survey the field. Our prune-during arm is gradual magnitude pruning; our contribution is the controlled comparison, the extreme-sparsity regime, and the local-rule variants.

**Dynamic sparse training.** SET [19] and RigL [7] start sparse and periodically drop and regrow connections (randomly or by gradient magnitude), with the Erdős–Rényi-kernel (ERK) layer allocation (we use ERK on the convolutional networks and a uniform allocation on the MLP). They are the cheapest of the large-network routes on the MLP and CNN and, at the tightest MNIST and ResNet-18 budgets, the strongest alternative to prune-during.

**Sparse attention and token skipping.** Sparse Transformers [2], top- $k$  attention [11] and routing attention [21] sparsify attention; DynamicViT [20] and A-ViT [27] drop tokens during a forward pass. Predictive coding [9, 18] motivates processing only prediction errors.

**Complementary learning systems.** The hippocampus–cortex account [15, 17] underlies replay-based continual learning: generative replay [23], CLS-ER [1], and class-balanced reservoir sampling [3]. Two implementations of the hippocampus–cortex structure do succeed at class-incremental learning without storing data: brain-inspired replay [26], in which the network’s own context-modulated feedback connections generate internal representations for replay, and CH-HNN [22], a hybrid of artificial and spiking networks that emulates the dual (specific and generalised) representations of corticohippocampal circuits and uses episode inference from prior knowledge. Our implementation in Section 5 is deliberately simpler—two plain MLPs, replay from a buffer, distillation—and its failure should be read against these successes. Synaptic homeostasis during sleep [25] motivated our downscaling variant. Pruning-based continual learning [10, 16] is related to the follow-up we propose in Section 6.

**STDP.** Our spiking network follows Diehl and Cook [4]. The accuracy ceiling of unsupervised STDP on shallow networks and its difficulty in scaling to deeper architectures are known [6]; the strongest STDP-only results add adaptive thresholds, adaptive synaptic filters and adaptive lateral inhibition [5]. We keep the plain rule because our interest is in what it does to the weight distribution, not in its accuracy.

### 3 Setup

**Routes to a sparse network.** Let  $B$  be a final budget of non-zero weights. We compare: *dense small*: a network of the same architecture whose width is scaled so that it has  $B$  weights, trained normally (additive); *prune-after*: the full-width network trained to completion, pruned once to  $B$  by global magnitude, then fine-tuned for half the original schedule; *prune-during*: the full-width network trained while a cubic schedule [28] lowers the density from 1 to  $B/|W|$  between 10% and 70% of training, pruning every 100 steps with no regrowth; *SET* and *RigL*: a random mask at density  $B/|W|$  (uniform per layer on the MLP, ERK on the CNN) with drop-and-regrow every 100 steps up to 75% of training, with random or gradient-based regrowth; *static sparse*: the random mask kept fixed. A full-width *dense big* network trained without pruning gives the accuracy ceiling.

**Pruning criteria.** Within *prune-during* we vary what is removed: weight magnitude  $|w_{ij}|$  (global or per-layer threshold); a Hebbian correlation trace  $|\mathbb{E}[y_i x_j]|$  with post-activation  $y_i$  and input  $x_j$ ; its product with magnitude; a *synaptic drive* trace  $|w_{ij}| \mathbb{E}[\mathbb{I}(y_i > 0) |x_j|]$ , i.e. the mean input carried by the synapse while its post-synaptic unit is active, which transports the “pre fires and post fires” condition of STDP to a synapse-level score without gradients; a per-neuron-normalised drive (dividing by the neuron’s mean drive, a homeostatic correction); and random removal as a control. All traces are exponential moving averages (momentum 0.99) maintained by forward hooks; convolutional drives are computed exactly per weight via unfolded receptive fields.

**Models and data.** MNIST: MLP 784–1024–1024–10 (1,861,632 weights), Adam, cosine schedule, 15 epochs, budgets 10%, 5%, 2%, 1% and 0.5%. CIFAR-10: VGG-style CNN with BatchNorm (conv 64-64/128-128/256-256, fc 256; 2,195,648 weights), SGD with Nesterov momentum, cosine schedule, 20 epochs, random crop and flip, budgets 10%, 3% and 1%. On the CNN a uniform per-layer density starves the first convolution (1,728 weights), so per-layer criteria use the ERK allocation as their layer targets. Three seeds for the MNIST, CIFAR-10 CNN and ResNet-18 comparisons; two for the transfer experiment, the inference-time pruning arms, the 20% ResNet-18 rows and the uniform-allocation ResNet-18 control; one for the CIFAR-10 ViT runs and for the 100- and 1,600-neuron spiking runs, which we mark as such. We report mean  $\pm$  standard deviation.

**Metrics.** Test accuracy; inference FLOPs counted analytically as  $2 \sum_{\ell} n_{\ell} d_{\ell}$  over linear and convolutional layers with  $n_{\ell}$  active weights and  $d_{\ell}$  output positions; cumulative training FLOPs as three times the forward FLOPs of the *actual* active network at every step; expected calibration error; data efficiency as the number of training samples needed to reach a target accuracy; and, for the sequential-task experiment, average final accuracy and forgetting. We do not measure energy: all experiments run on von Neumann GPUs (a consumer RTX 3060 Ti with 8 GB; the ResNet-18, transfer and inference-time pruning sweeps on one rented A40), where the brain’s energy argument cannot be tested.

### 4 Prune-during versus the alternatives

**Accuracy at matched budgets.** Tables 1 and 2 give the main result. Prune-during with a global magnitude criterion is the most accurate route at every budget of 2% and below. At the loosest budgets *prune-after* is level with it or slightly ahead: by 0.02 points at 10% and 0.01 at 5% on MNIST, 0.4 at 10% on CIFAR-10 and 0.3 at 5% on ResNet-18 (Table 3). Its margin over the additive baseline grows as the budget shrinks: on MNIST from 0.3 points at 10% to 3.7 points at

Table 1: MNIST test accuracy (%) at matched final connection budgets, 3 seeds. Standard deviations are below 0.1 for the magnitude routes at every budget and below 0.3 for every route at 2% and above; at 1% random removal and the static mask reach 0.6 and 0.3, and at 0.5% they reach 2.3 and 2.7, with prune-after 0.6, activity 1.1, RigL 0.3 and drive 0.2; full values in Table 7. Density is relative to the 1.86M-weight network. The dense big network reaches 98.7%.

| Route (criterion)                             | 10%          | 5%           | 2%           | 1%           | 0.5%         |
|-----------------------------------------------|--------------|--------------|--------------|--------------|--------------|
| Dense small (additive)                        | 98.36        | 98.01        | 97.30        | 96.30        | 94.07        |
| Prune-after + fine-tune                       | 98.69        | 98.73        | 98.38        | 97.52        | 94.51        |
| Prune-during, magnitude (global)              | <b>98.67</b> | <b>98.72</b> | <b>98.65</b> | <b>98.44</b> | <b>97.81</b> |
| Prune-during, magnitude (layer)               | 98.67        | 98.68        | 98.55        | 98.23        | 97.04        |
| Prune-during, synaptic drive                  | 98.53        | 98.50        | 98.22        | 97.56        | 96.03        |
| Prune-during, $ w  \times$ activity (Hebbian) | 98.57        | 98.42        | 98.05        | 97.25        | 95.22        |
| Prune-during, activity (Hebbian)              | 98.19        | 97.87        | 96.55        | 93.66        | 88.85        |
| Prune-during, random                          | 98.23        | 97.13        | 93.39        | 89.13        | 71.84        |
| RigL                                          | 98.53        | 98.22        | 97.85        | 97.41        | 96.31        |
| SET                                           | 98.38        | 98.19        | 97.43        | 96.03        | 93.65        |
| Static random sparse                          | 98.18        | 97.54        | 95.90        | 92.59        | 72.71        |

Table 2: CIFAR-10 test accuracy (%) and inference FLOPs per image at matched budgets (3 seeds). Density is relative to the 2.2M-weight CNN, which reaches 91.3% dense. Per-layer criteria use the ERK allocation. RigL at 1%: one of three seeds diverged to chance despite the non-finite-gradient guard (Appendix); the table shows the mean of the two seeds that converged, and the three-seed mean is  $53.9 \pm 31.1$  (Table 8, Figure 2).

| Route (criterion)                                | Accuracy    |             |                   | Inference FLOPs |       |       |
|--------------------------------------------------|-------------|-------------|-------------------|-----------------|-------|-------|
|                                                  | 10%         | 3%          | 1%                | 10%             | 3%    | 1%    |
| Dense small (additive)                           | 87.3        | 82.8        | 77.1              | 3.1e7           | 9.7e6 | 3.3e6 |
| Prune-after + fine-tune                          | <b>90.7</b> | 88.4        | 80.5              | 8.2e7           | 3.6e7 | 1.6e7 |
| Prune-during, magnitude (global)                 | 90.3        | <b>89.1</b> | <b>85.7</b>       | 7.4e7           | 3.0e7 | 1.2e7 |
| Prune-during, magnitude (ERK)                    | 89.3        | 86.9        | 81.3              | 2.7e7           | 8.4e6 | 2.8e6 |
| Prune-during, drive, per-neuron normalised (ERK) | 87.6        | 79.3        | 60.0              | 2.7e7           | 8.4e6 | 2.8e6 |
| RigL (ERK)                                       | 87.6        | 83.5        | 75.9 <sup>†</sup> | 2.7e7           | 8.4e6 | 2.8e6 |

0.5%, on CIFAR-10 from 3.0 points at 10% to 8.6 points at 1%. It also beats RigL by 1.5 (MNIST, 0.5%) and 9.8 points (CIFAR-10, 1%, against the two RigL seeds that converged), and prune-after by 3.3 and 5.2 points.

**The benefit is in the gradual removal.** Figure 1 shows why prune-after loses at high sparsity: the one-shot removal destroys the function and fine-tuning recovers only part of it. Gradual removal interleaved with learning never leaves the high-accuracy region. The same picture holds on CIFAR-10 at 1% (91.4% before the pruning step, 10% after it, 80.5% after fine-tuning; run logs). Sparsity itself is not the point: on MNIST the two routes that impose a random pattern, a static random mask and random removal during training, are the worst at every budget below 10% (72.7% and 71.8% at 0.5%).

**Connection counts versus FLOPs.** On the CNN, global magnitude pruning keeps convolutional weights, which are each used at many positions, and removes fully connected ones, so the pruned




Figure 1: Connection count (left, log scale) and test accuracy (right) during training at 0.5% density on MNIST (seed 0). Prune-after (orange) trains dense, is pruned once at step 7035, collapses below 60% and recovers to 94% during fine-tuning; the prune-during routes reach the same final count gradually; the global route (yellow) never falls below 97% (lowest point in any seed 97.3%, run logs) and the per-layer route (green) ends at 97.0%.

network uses  $2.4\text{--}3.5\times$  the inference FLOPs of the dense small network at the same connection count (Table 2). Constraining the per-layer allocation to ERK, which on this CNN gives *fewer* inference FLOPs than the dense small network, still leaves prune-during 4.2 points ahead at 1%, and 1.7–3.4 points ahead of RigL at the same FLOPs at 10% and 3% (5.4 at 1% against the converged RigL seeds; Figure 2). On ResNet-18 the ERK variant costs  $2.3\text{--}2.5\times$  the inference FLOPs of the dense small network, so there the equal-FLOPs comparison holds only against RigL (Table 3). The gain is therefore real in both currencies, but a paper that reports only connection counts would overstate it by a factor of two to three in compute.

**Training cost, data efficiency, calibration.** Prune-during starts with the full network, so its cumulative training FLOPs are  $3\times$  (10%) to  $50\times$  (0.5%) those of the dense small network on MNIST,  $4.5\text{--}31\times$  on CIFAR-10 and  $8\text{--}59\times$  on ResNet-18; prune-after costs a further  $2.5\text{--}4\times$ . RigL and SET cost the same as the dense small network on the MLP and CNN, and RigL  $2.3\text{--}2.5\times$  on ResNet-18. On the training-cost axis RigL is the most efficient of the large-network routes; in accuracy it is second to prune-during only at the tightest budgets (0.5% on MNIST and ResNet-18), otherwise third behind prune-after, and at 1% on CIFAR-10 it falls below the dense small network. Measured in samples rather than FLOPs, the over-parameterised network learns fastest, but on MNIST it reaches 97% after 87k samples, before pruning begins, so this is a property of the dense start that every large-network route shares (against 133k–436k samples for the dense small networks and 124k–445k for RigL). The comparison is informative where the target falls inside the pruning phase: on ResNet-18 at 5% prune-during reaches 85% after 385k samples, when about a fifth of the weights remain, earlier than the unpruned dense network (437k; prune-after is still dense at that point), RigL (524k) and the dense small network (665k); at 0.5% it needs 402k, against 437k, 908k (one RigL seed of three, the others never) and never. On MNIST the expected calibration error stays between 0.003 and 0.017 for the magnitude, additive, prune-after and RigL routes at every budget and degrades for the local rules at the tightest budgets (drive 0.025,  $|w|\times$ activity 0.038, activity 0.100), for the static random mask (0.053) and for random removal (0.157); on the CNN




Figure 2: CIFAR-10 accuracy against inference FLOPs; each point is one budget. At equal FLOPs, prune-during with the ERK allocation (orange) dominates RigL (green) and the dense small network (blue).

Table 3: CIFAR-10 ResNet-18 (11.2M weights, 92.5% dense): test accuracy (%) and inference FLOPs per image at matched budgets, three seeds. Per-layer criteria use the ERK allocation.

| Route (criterion)                | Accuracy    |             |             | Inference FLOPs |       |       |
|----------------------------------|-------------|-------------|-------------|-----------------|-------|-------|
|                                  | 5%          | 2%          | 0.5%        | 5%              | 2%    | 0.5%  |
| Dense small (additive)           | 88.7        | 86.5        | 81.3        | 5.6e7           | 2.3e7 | 6.1e6 |
| Prune-after + fine-tune          | <b>92.3</b> | 90.9        | 83.6        | 2.7e8           | 1.5e8 | 5.3e7 |
| Prune-during, magnitude (global) | 92.0        | <b>91.5</b> | <b>89.0</b> | 1.9e8           | 1.1e8 | 3.4e7 |
| Prune-during, magnitude (ERK)    | 92.0        | 91.3        | 87.8        | 1.4e8           | 5.7e7 | 1.4e7 |
| RigL (ERK)                       | 91.0        | 89.5        | 84.9        | 1.4e8           | 5.7e7 | 1.4e7 |

and ResNet-18 all routes lie between 0.008 and 0.038.

**ResNet-18.** Table 3 repeats the comparison with a CIFAR-style ResNet-18 (11.2M prunable weights, 92.5% dense; SGD, 20 epochs, three seeds) at 5%, 2% and 0.5% of the weights. The winner is unchanged: at 0.5% (56k weights) prune-during reaches 89.0% against 81.3% for the dense small network, 83.6% for prune-after and 84.9% for RigL, its largest lead over prune-after in any setting (5.3 points) and a lead over the additive baseline (7.7) close to the CNN’s (8.6); constrained to the ERK allocation it keeps 87.8% at 40% of the inference FLOPs of the global variant. The order of the other routes does change with the model: on the CNN at 1% prune-after beats dense small beats RigL, on ResNet-18 at 0.5% RigL beats prune-after beats dense small. At 5% the three large-network routes lie within 0.3 points of each other and within 0.5 points of the dense ceiling, so the budget, not the method, decides whether the route matters.

**What to remove: local rules.** Removing connections by a Hebbian correlation trace beats random removal but falls below the dense small network at every budget under 10% (88.9% against 94.1% at 0.5%), because  $\mathbb{E}[y_i x_j]$  measures correlation through every path, not the contribution of synapse  $ij$ . Scoring synapses by their own drive repairs most of the gap on the MLP: at 0.5% the drive rule (96.0%) beats the additive baseline, prune-after and SET and is 0.3 points behind RigL, with a criterion that uses no gradient information (training itself remains back-propagation). On the CNN the unnormalised drive rule collapses within two epochs (run logs): channels that fire rarely lose all their inputs, die, and the failure propagates downstream, a known hazard of activity-dependent pruning without homeostasis. Normalising each neuron’s drive by its mean removes the collapse but the rule still falls from parity with the dense baseline at 10% to 17 points behind it at 1%. A local criterion can stand in for the magnitude criterion on the MLP at a cost of 1.0 point at 0.5% against the per-layer magnitude criterion (1.8 against the global one); on the convolutional network, in this form, it cannot.

**Transfer: pruning a pre-trained network during adaptation.** The over-parameterised starting point that prune-during needs is, in practice, often inherited: pre-trained networks are published with their training cost already paid, so the practical question is what gradual pruning does during adaptation. We therefore took an ImageNet-pre-trained ResNet-18 (torchvision weights, 11.17M convolutional weights, fresh 10-way classifier kept dense) and adapted it to CIFAR-10 for 10 epochs at  $128 \times 128$  input while pruning its convolutions to 5%, 2% and 0.5% with the same cubic schedule (SGD, lr 0.01, the from-scratch arms use lr 0.1 and the same architecture, schedule and resolution). Table 4 compares gradual pruning during adaptation with one-shot magnitude pruning of the same weights followed by fine-tuning, RigL started from a random ERK mask over the pre-trained weights, and the two from-scratch routes. At 5% and 2% the pre-trained start is worth 5.3 and 3.5 points over pruning from scratch, gradual removal during adaptation beats one-shot pruning of the same weights by 0.6–0.7 points, and RigL, which must discard the inherited structure to start sparse, lands below pruning from scratch. The pre-trained prune-during arms reach 90% after about 52k samples (just over one epoch, when pruning has just begun and about 99% of the weights remain, as does the dense fine-tune), a level one-shot pruning needs 166k–402k samples for at 5% and 2% and does not reach at 0.5%, and RigL never reaches in ten epochs; they are also better calibrated (ECE 0.015–0.020 against 0.028–0.031 for prune-during from scratch). The inherited start does not save adaptation compute: at 5% gradual pruning costs 7.3e14 FLOPs from the pre-trained weights against 7.7e14 from scratch, 40% of a dense fine-tune and  $7 \times$  the dense small network. What it buys is accuracy; the early data efficiency is that of the dense start, which gradual pruning keeps and one-shot pruning destroys. At 0.5% the ordering reverses: pruning from scratch (83.8%) beats one-shot pruning (81.7%), which beats gradual pruning of the pre-trained network (80.9%). The pre-trained prune-during run peaks at 93.1% in mid-schedule and loses twelve points as the last 4.5% of its connections go, with only three epochs left at the final density, whereas one-shot pruning spends all ten there. An inherited structure pays off down to about 2% of its connections; below that, a magnitude profile tuned for another task is a liability, and the short adaptation window favours removing it at once.

## 5 Negative results from the same programme

The programme that produced Section 4 also implemented the other mechanisms of Table 5. Four of them did not work as designed, and we report them because the failures are informative: sleep consolidation, thalamic routing, predictive coding and per-input subnetworks. The local pruning

Table 4: ImageNet-pre-trained ResNet-18 adapted to CIFAR-10 ( $128 \times 128$ , 10 epochs) while being pruned: test accuracy (%) at matched convolutional budgets, two seeds, standard deviations  $\leq 0.4$ . Adaptation FLOPs count the 10 epochs only; the ImageNet pre-training is a sunk cost common to the pre-trained arms. The unpruned dense fine-tune reaches 95.9% at  $1.8 \times 10^{15}$  FLOPs.

| Start       | Route                            | 5%          | 2%          | 0.5%        | Adapt. FLOPs (5%) |
|-------------|----------------------------------|-------------|-------------|-------------|-------------------|
| pre-trained | prune-during, magnitude (global) | <b>93.7</b> | <b>91.0</b> | 80.9        | 7.3e14            |
| pre-trained | prune-during, magnitude (ERK)    | 93.4        | 89.7        | 78.8        | 8.1e14            |
| pre-trained | one-shot prune + fine-tune       | 93.1        | 90.3        | 81.7        | 3.0e14            |
| pre-trained | random ERK mask + RigL           | 86.9        | 83.6        | 76.8        | 2.7e14            |
| scratch     | prune-during, magnitude (global) | 88.4        | 87.5        | <b>83.8</b> | 7.7e14            |
| scratch     | dense small (additive)           | 84.8        | 82.4        | 76.5        | 1.1e14            |




Figure 3: Pre-trained ResNet-18 adapted to CIFAR-10 while pruned: accuracy against final active convolutional weights (two seeds). Solid lines start from the ImageNet weights, dashed lines from scratch. The pre-trained advantage is largest at 5% and reverses at 0.5%.

criteria of Section 4 are a variant of the first mechanism and are not counted separately; the spiking network closes the section as a replication, not a failure.

**Sleep consolidation.** We trained a fast “hippocampal” MLP on the task stream and a slow “cortical” MLP only during sleep phases at task boundaries, using the current task’s data plus a 500-sample reservoir buffer, with distillation from the fast network, global synaptic downscaling and magnitude or activity pruning as optional sleep ingredients (Split MNIST class-incremental, Permuted MNIST, MLP 784–256–256–10, three seeds; Tables 11 and 12). Distillation from the fast network lowered final average accuracy on Split MNIST from 85.0% to 78.7% and raised forgetting from 0.18 to 0.25, because the fast network, trained on the current task’s classes alone under the same single 10-way head, is biased to them; downscaling hurt in both settings (77.4% and 82.7%); magnitude pruning was neutral and activity pruning cost 6.3 points on Permuted MNIST; and the scheme without distillation only matched plain experience replay (85.0% against 85.2%) at three

Table 5: The mechanisms ported in this paper, their reductions, and what transferred.

| Brain mechanism                                | Our reduction                                             | Outcome                                                                                                       | Where        |
|------------------------------------------------|-----------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|--------------|
| gradual synaptic pruning during learning       | prune-during at matched budgets                           | transfers (principle)                                                                                         | Section 4    |
| activity-dependent local pruning rules         | Hebbian, drive and normalised-drive criteria              | does not transfer: trails magnitude pruning, collapses on the CNN without homeostasis, near parity on the MLP | Section 4    |
| hippocampus-cortex sleep consolidation         | fast/slow MLPs, replay, distillation, decay               | matches class-balanced replay; beats plain replay by 1.2 points on Permuted MNIST at equal compute            | this section |
| thalamic routing                               | pre-decided skips; pre-routed attention                   | does not transfer (choice uninformative; little compute saved)                                                | this section |
| predictive coding                              | skip by predicted change; predicted-residual substitution | does not transfer: selection no better than random, substitution no better than identity                      | this section |
| per-input subnetworks at inference (follow-up) | per-sample masks by a local rule                          | does not transfer                                                                                             | this section |
| spike-timing-dependent plasticity              | Diehl-Cook network                                        | known accuracy limit reproduced; self-sparsification observed                                                 | this section |

times the compute. A redesign with a  $k$ -winners-take-all sparse hippocampus, no distillation and class-balanced replay improved to 87.0%, but balanced sampling alone added to experience replay reaches 87.8%, and the hippocampal sparsity has no path to influence the cortex once distillation is removed. On Split MNIST the only effective ingredient was the known one [3]; on Permuted MNIST the scheme without distillation beats experience replay by 1.2 points at the same compute (88.6% against 87.4%), and the sparse redesign beats balanced replay by 0.8 points at five times the compute (89.0% against 88.3%). This is a negative result for our simple implementation, not for the hippocampus-cortex account: brain-inspired generative replay [26] and the hybrid CH-HNN [22] both succeed at class-incremental learning, and what they have that we lacked is a hippocampus with generative capacity (internal representations replayed through feedback connections) or a dual specific-generalised representation with episode inference, rather than a second feed-forward network that can only speak to the cortex through distillation. Confidence-based agreement between the two networks, proposed as a metacognitive “I do not know” signal, failed to detect unseen classes (AUROC 0.51), whereas plain softmax confidence reached 0.85. Varying how the sleep phase decays weights, on the sparse redesign (Tables 19 and 20), did not change the picture: a one-shot 10% downscaling at the task boundary, a gradual 10% or 3% decay spread over the sleep phase and a continuous weight decay were within noise on Split MNIST (+0.1 to +0.5 points) and cost 0.3–4.0 points on Permuted MNIST. Replacing global logit distillation by layer-local feature distillation, the form a processor-memory split would favour, did not help: with a 256-unit fast network, local distillation costs 4.2 points against no distillation on both benchmarks, less than global distillation on Split MNIST (6.2) and more on Permuted MNIST (1.6).

**Predictive token skipping and pre-filtering.** With a low-rank predictor per block estimating how much each token will change, skipping the tokens predicted to change least, uniformly over all blocks of a small ViT, was *worse* than skipping random tokens (MNIST, 30% skip: 94.9% against 96.7%; even the oracle residual norm was worse than random), because background tokens freeze consistently in every layer. Restricting skipping to the later half of the blocks recovers most of the loss by itself: with the skipped tokens left unchanged and no fine-tuning, skipping 70% of them




Figure 4: CIFAR-10 ViT: accuracy against FLOPs when skipping tokens in the late blocks with predicted-residual substitution, after fine-tuning with a ramped skip ratio (one seed). Thalamic pre-decision, per-block prediction, random selection and identity substitution (no predicted residual) coincide; all lie far above post-hoc identity skipping without fine-tuning (grey).

in the late blocks costs 0.3 points on MNIST and 5.0 on CIFAR-10, against 1.5–3.3 and 8.9–15.4 points for skipping 30% of them in every block at the same FLOPs (Tables 13 and 14), and a short fine-tuning phase that ramps the skip ratio from 0 to 0.7 brings the CIFAR-10 cost to 1.0 point at 72% of the FLOPs (ViT, 81.8% baseline, one seed; Figure 4). Replacing the skipped tokens by their predicted residual, the predictive-coding ingredient, adds nothing over leaving them unchanged under the same fine-tuning: 97.9% against 97.9% on MNIST (three seeds) and 80.8% against 80.8% on CIFAR-10 at 70% skipping (one seed; within 0.7 points at every non-zero ratio, with a 1.0-point gap at zero skipping that sets the single-seed noise floor; Tables 15 and 16). A “thalamic” router that pre-decides all late-block skips from the representation at the block boundary performs identically to per-block prediction and to random selection, also under cross-evaluation (three seeds on MNIST, one on CIFAR-10; Tables 15 and 16). The choice of *which* tokens to skip carries no information at this scale, and neither does what replaces them; what matters is how many and in which layers.

**Pre-routed sparse attention.** Selecting the top- $k$  keys per query with an 8-dimensional router *before* the full attention computation keeps accuracy at  $k = 5$  of 50 keys on MNIST (98.1% against 97.9%, three seeds) and  $k = 17$  of 65 on CIFAR-10 (80.0% against 79.4%, one seed; Tables 17 and 18) and cuts attention FLOPs by 68–77%, but attention is only 5–6% of the total cost at 50–65 tokens, so total FLOPs fall by only 3.6–4.7%; post-hoc top- $k$  masking still computes the full score matrix and saves at most 2.7%. Stacked on token skipping it saves a further 1% of FLOPs for a 0.1–0.6 point loss. The mechanism is sound but irrelevant below several hundred tokens.

**Input-dependent pruning at inference time.** If learning is pruning, should inference prune as well? We let the ResNet-18 of Table 3 select a different subnetwork for every input. For each output neuron, the connections with the largest  $|w_{ij}| \overline{|x_j|}$  (the weight times the mean magnitude of its input

Table 6: Input-dependent (per-sample) pruning of a CIFAR-10 ResNet-18 against fixed masks at the same per-sample budget. Static rows from Table 3 (three seeds; the 20% column and the uniform-allocation control from Table 9, two seeds); dynamic rows two seeds, with per-sample budgets within 1% of the fixed masks at 20% and 5% and 14% above them at 0.5%. Jaccard: overlap of the last-stage masks of two test images of the same class / of different classes (200 images, 5% budget); a random mask gives 0.026.

| Mask                                               | 20%         | 5%          | 0.5%        | Jaccard same / diff. class (5%) |
|----------------------------------------------------|-------------|-------------|-------------|---------------------------------|
| fixed, prune-during, magnitude (global)            | <b>91.9</b> | <b>92.0</b> | <b>89.0</b> | –                               |
| fixed, prune-during, magnitude (uniform per layer) | 91.8        | 91.1        | 83.1        | –                               |
| fixed, dense small (additive)                      | 90.9        | 88.7        | 81.3        | –                               |
| per sample, connection-level, local rule           | 91.8        | 89.8        | 77.7        | 0.44 / 0.27                     |
| per sample, connection-level, random               | 76.1        | 52.3        | 18.2        | 0.03 / 0.03                     |
| per sample, channel-level $k$ -WTA gate            | 73.4        | 44.9        | 34.0        | 0.08 / 0.05                     |

channel for that sample, a per-sample form of the synaptic-drive rule without its post-synaptic term) are kept and the rest are masked for that sample only, with the kept fraction following the same cubic schedule as prune-during so that every sample uses the budget (within 1% at 20% and 5%; 14% more at 0.5%, where per-neuron rounding and the dense first layer dominate). Controls: the same per-sample budget chosen at random, a channel-level  $k$ -WTA gate on each convolution’s input (the only variant whose sparsity is structured enough to save time on a GPU), and the static networks of Table 3 plus a static mask with the same uniform per-layer allocation that the dynamic rule is confined to. Table 6 gives the result. At 20% the input-dependent rule matches the fixed masks (91.8% against 91.8–91.9%). The penalty appears as the budget shrinks: at 5% it reaches 89.8% against 92.0% for a fixed global mask and 91.1% for a fixed mask with the same per-layer allocation, and at 0.5% 77.7% against 89.0% and 83.1%, so re-deciding the mask for every input costs 1.3 points at 5% and 5.4 points at 0.5% even after allocation is controlled for; it beats the width-matched dense network at 20% and 5% (90.9%, 88.7%) and loses to it at 0.5% (81.3%). In FLOPs the dynamic arm sits between the two fixed masks (at 0.5%, 9.3e6 against 5.6e6 for the uniform and 3.4e7 for the global mask). Random per-input selection collapses to 52.3% at 5% and 18.2% at 0.5%, and the channel-level gate to 44.9% and 34.0% (73.4% even at 20%). The rule does form input-specific subnetworks: in the last stage, the masks of two samples of the same class share 44% of their connections at 5% and those of different classes 27% (random masks: 2.6%; at 20%, 61% against 46%), and across 200 test images the network touches 21% of that stage’s connections while each image uses 5%. The extra capacity is real and buys nothing. This is the mirror image of the token result: among tokens, which to drop did not matter; among connections, which to keep is everything, and re-deciding it for every input is worse than deciding it once.

**STDP sparsifies itself.** A 400-neuron spiking network in the style of Diehl and Cook [4], trained with trace-based STDP, lateral inhibition and weight normalisation on one pass of MNIST, reaches  $81.1 \pm 0.3\%$  (a back-propagated MLP of the same width reaches 96.5% on the same samples). The gap is the known limitation of unsupervised STDP [5, 6]; what we add is an observation about the weights. Without any explicit pruning, the fraction of its input weights below 1% of the maximum rises from 5.0% to 57.7% and the Gini coefficient of the weight distribution from 0.33 to 0.77 during learning (run logs; the 100- and 1,600-neuron runs end at 54.9% and 28.7%, Table 22). Learning under a local plasticity rule with normalisation is itself subtractive, which is the biological reading of Section 4.

## 6 Discussion and limitations

**Scope of the claim.** The claim that survives is narrow and, we think, useful: at a tight final connection budget, removing connections gradually while learning is better than learning small, better than removing them afterwards and better than rewiring from a sparse start; the advantage grows with sparsity, and at loose budgets it shrinks to a tie with removing them afterwards. The claim is about accuracy per connection. It is not a claim about training cost, where RigL wins, and on convolutional networks it must be stated per FLOP as well as per connection. It is also not a claim that local rules suffice: they nearly do on the MLP and do not, in our forms, on the CNN.

**Training cost, and what an inherited start buys.** The  $3\text{--}59\times$  training-cost multiplier is the honest weakness of prune-during trained from scratch: it buys accuracy with compute, and RigL buys most of the accuracy for the price of the dense small network ( $2.3\text{--}2.5\times$  on ResNet-18), trailing prune-during by 1.5–10 points at the tightest budgets. Two observations qualify the weakness. First, training happens once and inference many times; at equal inference FLOPs (the ERK-constrained variant) the extra training compute is amortised over deployment, so the trade-off depends on how often the model is used. Second, the over-parameterised starting point is often inherited: pre-trained networks are published with their training cost already paid. That does not make the adaptation cheap, since the schedule still runs its dense phase (at 5%,  $7.3e14$  FLOPs against  $7.7e14$  from scratch), but it changes what the compute buys. Pruning pre-trained models is the classical setting of Han et al. [12] and Zhu and Gupta [28]; what our comparison adds is that gradual removal *during* adaptation should then dominate one-shot pruning and, because RigL must discard the pre-trained structure to start sparse, dynamic sparse training as well. The transfer experiment (Table 4) confirms this at 5% and 2%: gradual removal during adaptation beats one-shot pruning of the same weights, beats RigL on the pre-trained weights by seven points and keeps 5.3 and 3.5 points over pruning from scratch, at 40% of the FLOPs of a dense fine-tune. It refutes it at 0.5%, where the inherited structure becomes a liability and both one-shot pruning and training from scratch win: an inherited start pays off only while enough of it survives. The developmental reading is the same: an organism inherits an over-connected cortex it did not pay for and only pays for the pruning.

**Scale.** Two small datasets; models from one to eleven million parameters. The ResNet-18 replication (Table 3) shows that the winner and the growth of the margin with sparsity survive a five-fold increase in model size, but ImageNet-scale data and models remain future work. The hazards we found on the CNN (layer starvation under uniform allocation, channel death under activity pruning, divergence of RigL at sub-1% mid-layer densities) suggest that the extreme-sparsity regime will need care at scale.

**What the negative results constrain.** A hippocampus that cannot communicate with the cortex except by distillation will transfer its bias; a token-selection rule that freezes the same background tokens in every layer will lose to random selection; an attention-sparsity mechanism is only worth its router at long sequences; a network that re-selects its connections for every input pays for the flexibility with accuracy. These are not arguments against the brain principles but against the particular reductions we tried, and they are the kind of result that is rarely published and often re-discovered.

**An observation about the failures.** We state the distinction before the results. We call a reduction a *structure* when it adds a module or a decision to the network (a second network, a router, a predictor that fills gaps, a per-sample mask) and a *principle* when it changes only how the existing network is trained (what is removed, and when). By this criterion the four mechanisms that did not transfer are all structures, and the one that did is the only principle we tested; the local pruning rules sit in between (a different criterion for the same schedule, with activity buffers added) and transferred only in part. One reading, which we offer as a hypothesis rather than a conclusion because it was named after the results and rests on a single positive case, is that a substantial part of the brain’s advantage rests on compute–memory unity, the hardware property by which every synapse computes, stores and learns in place, and that a processor–memory architecture can only imitate the structures at a cost: a hippocampus that can communicate only by distillation transfers its bias, a router must spend compute to save compute, a predictor must be trained and evaluated for every token it stands in for, an activity-dependent rule must track activity in buffers separate from the weights it governs, and a per-input subnetwork must be materialised as a per-sample copy of the weights. Each failure also has a mundane explanation (bias transfer, a single shared head in class-incremental learning, attention being 5–6% of the cost, token selection carrying no information at this scale, a predicted residual carrying no more than the token it replaces once only late layers are skipped, a mask that changes faster than the weights can adapt to it), and the hippocampus–cortex structure has been made to work in software by giving the hippocampus generative capacity or a hybrid dual representation [22, 26]. We did not test the hardware reading: the principle transferred; the structures did not, in our hands.

**Follow-up.** The natural next question links the main result to forgetting: whether networks sparsified by prune-during interfere less across sequential tasks than dense networks of equal budget. This connects to PackNet-style continual learning [10, 16] but tests the developmental claim directly.

## 7 Conclusion

Pruning during learning, with a plain global magnitude criterion and a gradual schedule, is the best route to a sparse network at tight budgets on the three models we tried, ties with pruning afterwards at loose budgets, and its advantage comes from the gradualness rather than from sparsity. Starting from a pre-trained network adds 3.5–5.3 points over pruning from scratch down to about 2% of the connections, without saving adaptation compute, and becomes a liability below that. The same programme’s attempts at sleep consolidation, predictive skipping, thalamic routing (pre-routed attention) and input-dependent pruning did not add to known baselines, apart from a 1.2-point gain over plain replay on Permuted MNIST at equal compute, and we report them so that the next attempt can start from the right place.

**Reproducibility.** Code, configurations, per-run logs and the full result tables will be released in a public repository; every number in this paper is produced by `scripts/analyze.py` from the stored run files.

## References

- [1] Elahe Arani, Fahad Sarfraz, and Bahram Zonooz. Learning fast, learning slow: a general continual learning method based on complementary learning system. In *International Conference on Learning Representations*, 2022.

- [2] Rewon Child, Scott Gray, Alec Radford, and Ilya Sutskever. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019.
- [3] Aristotelis Chrysakis and Marie-Francine Moens. Online continual learning from imbalanced data. In *International Conference on Machine Learning*, 2020.
- [4] Peter U. Diehl and Matthew Cook. Unsupervised learning of digit recognition using spike-timing-dependent plasticity. *Frontiers in Computational Neuroscience*, 9:99, 2015.
- [5] Yiting Dong, Dongcheng Zhao, Yang Li, and Yi Zeng. An unsupervised STDP-based spiking neural network inspired by biologically plausible learning rules and connections. *Neural Networks*, 165:799–808, 2023. doi: 10.1016/j.neunet.2023.06.019.
- [6] Jason K. Eshraghian, Max Ward, Emre O. Neftci, Xinxin Wang, Gregor Lenz, Girish Dwivedi, Mohammed Bennamoun, Doo Seok Jeong, and Wei D. Lu. Training spiking neural networks using lessons from deep learning. *Proceedings of the IEEE*, 111(9):1016–1054, 2023. doi: 10.1109/JPROC.2023.3308088.
- [7] Utku Evci, Trevor Gale, Jacob Menick, Pablo Samuel Castro, and Erich Elsen. Rigging the lottery: making all tickets winners. In *International Conference on Machine Learning*, 2020.
- [8] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: finding sparse, trainable neural networks. In *International Conference on Learning Representations*, 2019.
- [9] Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11:127–138, 2010.
- [10] Siavash Golkar, Michael Kagan, and Kyunghyun Cho. Continual learning via neural pruning. *arXiv preprint arXiv:1903.04476*, 2019.
- [11] Ankit Gupta, Guy Dar, Shaya Goodman, David Ciprut, and Jonathan Berant. Memory-efficient transformers via top- $k$  attention. In *Proceedings of the Second Workshop on Simple and Efficient Natural Language Processing*, 2021.
- [12] Song Han, Jeff Pool, John Tran, and William J. Dally. Learning both weights and connections for efficient neural networks. In *Advances in Neural Information Processing Systems*, 2015.
- [13] Torsten Hoeftler, Dan Alistarh, Tal Ben-Nun, Nikoli Dryden, and Alexandra Peste. Sparsity in deep learning: pruning and growth for efficient inference and training in neural networks. *Journal of Machine Learning Research*, 22(241):1–124, 2021.
- [14] Peter R. Huttenlocher. Synaptic density in human frontal cortex: developmental changes and effects of aging. *Brain Research*, 163(2):195–205, 1979.
- [15] Dharshan Kumaran, Demis Hassabis, and James L. McClelland. What learning systems do intelligent agents need? Complementary learning systems theory updated. *Trends in Cognitive Sciences*, 20(7):512–534, 2016.
- [16] Arun Mallya and Svetlana Lazebnik. PackNet: adding multiple tasks to a single network by iterative pruning. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018.

- [17] James L. McClelland, Bruce L. McNaughton, and Randall C. O'Reilly. Why there are complementary learning systems in the hippocampus and neocortex: insights from the successes and failures of connectionist models of learning and memory. *Psychological Review*, 102(3): 419–457, 1995.
- [18] Beren Millidge, Alexander Tschantz, and Christopher L. Buckley. Predictive coding approximates backprop along arbitrary computation graphs. *Neural Computation*, 34(6):1329–1368, 2022.
- [19] Decebal Constantin Mocanu, Elena Mocanu, Peter Stone, Phuong H. Nguyen, Madeleine Gibescu, and Antonio Liotta. Scalable training of artificial neural networks with adaptive sparse connectivity inspired by network science. *Nature Communications*, 9:2383, 2018.
- [20] Yongming Rao, Wenliang Zhao, Benlin Liu, Jiwen Lu, Jie Zhou, and Cho-Jui Hsieh. DynamicViT: efficient vision transformers with dynamic token sparsification. In *Advances in Neural Information Processing Systems*, 2021.
- [21] Aurko Roy, Mohammad Saffar, Ashish Vaswani, and David Grangier. Efficient content-based sparse attention with routing transformers. *Transactions of the Association for Computational Linguistics*, 9:53–68, 2021.
- [22] Qianqian Shi, Faqiang Liu, Hongyi Li, Guangyu Li, Luping Shi, and Rong Zhao. Hybrid neural networks for continual learning inspired by corticohippocampal circuits. *Nature Communications*, 16:1272, 2025. doi: 10.1038/s41467-025-56405-9.
- [23] Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. Continual learning with deep generative replay. In *Advances in Neural Information Processing Systems*, 2017.
- [24] Hidenori Tanaka, Daniel Kunin, Daniel L. K. Yamins, and Surya Ganguli. Pruning neural networks without any data by iteratively conserving synaptic flow. In *Advances in Neural Information Processing Systems*, 2020.
- [25] Giulio Tononi and Chiara Cirelli. Sleep and the price of plasticity: from synaptic and cellular homeostasis to memory consolidation and integration. *Neuron*, 81(1):12–34, 2014.
- [26] Guido M. van de Ven, Hava T. Siegelmann, and Andreas S. Tolias. Brain-inspired replay for continual learning with artificial neural networks. *Nature Communications*, 11:4069, 2020. doi: 10.1038/s41467-020-17866-2.
- [27] Hongxu Yin, Arash Vahdat, Jose M. Alvarez, Arun Mallya, Jan Kautz, and Pavlo Molchanov. A-ViT: adaptive tokens for efficient vision transformer. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- [28] Michael Zhu and Suyog Gupta. To prune, or not to prune: exploring the efficacy of pruning for model compression. *arXiv preprint arXiv:1710.01878*, 2017.

## A Hyper-parameters and implementation notes

**Masks.** Linear and convolutional layers carry a binary mask buffer; the forward pass uses  $w \odot m$ , masked weights are re-zeroed after every optimiser step, and pruning selects an exact top- $k$  of the live weights by index so that tied scores (common with activity traces) cannot inflate the count.

**Schedules.** Prune-during: density  $d(t) = d_f + (1 - d_f)(1 - t)^3$  for  $t \in [0, 1]$  spanning 10–70% of training, pruning every 100 steps, with a final pruning step at 70% so that the target is met exactly. RigL/SET: drop fraction  $0.3 \cdot \frac{1}{2}(1 + \cos(\pi s/S))$  every 100 steps up to  $S = 75\%$  of training. Prune-after: fine-tuning for half the original epochs at lr 0.01 (CNN) or the original lr (MLP). RigL regrowth is skipped on steps whose mixed-precision gradients are non-finite; without this guard one CIFAR-10 seed lost 93% of its connections in the first 1,000 steps. With the guard in place one CIFAR-10 seed at 1% still diverged to chance (Table 2).

**Dense small sizing.** MLP: hidden width  $h$  solving  $784h + h^2 + 10h = B$ . CNN and ResNet: all widths scaled by a common factor found by bisection on the prunable weight count.

**Training.** MLP: Adam, lr  $10^{-3}$ , batch 128, one warm-up epoch then cosine, 15 epochs, no augmentation, full training and test sets resident on the GPU. CNN: SGD with Nesterov momentum 0.9, lr 0.05, weight decay  $5 \times 10^{-4}$  on weights only, batch 128, mixed precision, 20 epochs, random crop (padding 4) and horizontal flip performed on the GPU.

**Spiking network.** 784 Poisson inputs at 32 Hz per unit intensity, 250 ms per image, 1 ms steps; 400 LIF neurons with adaptive thresholds ( $\theta^+ = 0.05$ ), 5 ms refractory period, recurrent inhibition of 17.5 mV; trace STDP with  $\eta_{\text{pre}} = 10^{-4}$ ,  $\eta_{\text{post}} = 10^{-2}$ ; input weights normalised to a column sum of 78.4; batches of 32 images with summed updates. Labels are assigned to neurons by their class-wise mean response on 10k training images.

## B Full result tables

Tables 7–22 give every number behind the figures and tables of the main text, as mean  $\pm$  standard deviation over seeds. FLOPs are analytic counts for one sample (inference) or accumulated over training at the actual density of each step. Five statements in the text rest on run logs rather than on these tables: the two-epoch collapse of the unnormalised drive rule on the CNN, the accuracy trajectory of prune-after on CIFAR-10 at its pruning step, the lowest point of the prune-during trajectory in Figure 1, the initial weight statistics of the spiking network, and its final Gini coefficient.

| density | arm                                            | seeds | active weights | test acc            | ECE               | infer FLOPs        | cum. train FLOPs      | samples to 97%     |
|---------|------------------------------------------------|-------|----------------|---------------------|-------------------|--------------------|-----------------------|--------------------|
| 1       | dense big (no pruning)                         | 3     | 1,861,632      | 0.9871 $\pm$ 0.0003 | 0.009 $\pm$ 0.000 | $3.72 \times 10^6$ | $1.01 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.1     | dense small (additive)                         | 3     | 185,787        | 0.9836 $\pm$ 0.0003 | 0.009 $\pm$ 0.000 | $3.72 \times 10^5$ | $1.00 \times 10^{12}$ | 133096 $\pm$ 6988  |
| 0.1     | train-then-prune + finetune                    | 3     | 186,163        | 0.9869 $\pm$ 0.0003 | 0.008 $\pm$ 0.000 | $3.72 \times 10^5$ | $1.06 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: magnitude (layer)                | 3     | 186,164        | 0.9867 $\pm$ 0.0003 | 0.007 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: magnitude (global)               | 3     | 186,163        | 0.9867 $\pm$ 0.0002 | 0.008 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: activity (layer)                 | 3     | 186,164        | 0.9819 $\pm$ 0.0007 | 0.004 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: abs(w) $\times$ activity (layer) | 3     | 186,164        | 0.9857 $\pm$ 0.0004 | 0.006 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: synaptic drive (layer)           | 3     | 186,164        | 0.9853 $\pm$ 0.0004 | 0.007 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | prune-during: random (layer)                   | 3     | 186,164        | 0.9823 $\pm$ 0.0012 | 0.004 $\pm$ 0.000 | $3.72 \times 10^5$ | $3.33 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.1     | SET (dynamic, random regrow)                   | 3     | 186,164        | 0.9838 $\pm$ 0.0004 | 0.010 $\pm$ 0.001 | $3.72 \times 10^5$ | $1.01 \times 10^{12}$ | 138766 $\pm$ 2328  |
| 0.1     | RigL (dynamic, grad regrow)                    | 3     | 186,164        | 0.9853 $\pm$ 0.0011 | 0.008 $\pm$ 0.001 | $3.72 \times 10^5$ | $1.01 \times 10^{12}$ | 123519 $\pm$ 1928  |
| 0.1     | static random sparse                           | 3     | 186,164        | 0.9818 $\pm$ 0.0002 | 0.006 $\pm$ 0.000 | $3.72 \times 10^5$ | $1.01 \times 10^{12}$ | 206096 $\pm$ 4555  |
| 0.05    | dense small (additive)                         | 3     | 93,392         | 0.9801 $\pm$ 0.0003 | 0.009 $\pm$ 0.001 | $1.87 \times 10^5$ | $5.04 \times 10^{11}$ | 186445 $\pm$ 12401 |
| 0.05    | train-then-prune + finetune                    | 3     | 93,082         | 0.9873 $\pm$ 0.0003 | 0.006 $\pm$ 0.000 | $1.86 \times 10^5$ | $1.03 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: magnitude (layer)                | 3     | 93,082         | 0.9868 $\pm$ 0.0006 | 0.006 $\pm$ 0.001 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: magnitude (global)               | 3     | 93,082         | 0.9872 $\pm$ 0.0003 | 0.006 $\pm$ 0.001 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: activity (layer)                 | 3     | 93,082         | 0.9787 $\pm$ 0.0001 | 0.005 $\pm$ 0.001 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: abs(w) $\times$ activity (layer) | 3     | 93,082         | 0.9842 $\pm$ 0.0004 | 0.004 $\pm$ 0.000 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: synaptic drive (layer)           | 3     | 93,082         | 0.9850 $\pm$ 0.0003 | 0.005 $\pm$ 0.000 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | prune-during: random (layer)                   | 3     | 93,082         | 0.9713 $\pm$ 0.0011 | 0.015 $\pm$ 0.000 | $1.86 \times 10^5$ | $2.96 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.05    | SET (dynamic, random regrow)                   | 3     | 93,082         | 0.9819 $\pm$ 0.0009 | 0.006 $\pm$ 0.001 | $1.86 \times 10^5$ | $5.03 \times 10^{11}$ | 209140 $\pm$ 5774  |
| 0.05    | RigL (dynamic, grad regrow)                    | 3     | 93,082         | 0.9822 $\pm$ 0.0004 | 0.008 $\pm$ 0.000 | $1.86 \times 10^5$ | $5.03 \times 10^{11}$ | 153687 $\pm$ 5905  |
| 0.05    | static random sparse                           | 3     | 93,082         | 0.9754 $\pm$ 0.0003 | 0.004 $\pm$ 0.001 | $1.86 \times 10^5$ | $5.03 \times 10^{11}$ | 378278 $\pm$ 18120 |
| 0.02    | dense small (additive)                         | 3     | 36,872         | 0.9730 $\pm$ 0.0010 | 0.006 $\pm$ 0.000 | $7.37 \times 10^4$ | $1.99 \times 10^{11}$ | 436301 $\pm$ 46700 |
| 0.02    | train-then-prune + finetune                    | 3     | 37,233         | 0.9838 $\pm$ 0.0008 | 0.003 $\pm$ 0.001 | $7.45 \times 10^4$ | $1.02 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: magnitude (layer)                | 3     | 37,233         | 0.9855 $\pm$ 0.0007 | 0.003 $\pm$ 0.000 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: magnitude (global)               | 3     | 37,233         | 0.9865 $\pm$ 0.0002 | 0.003 $\pm$ 0.001 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: activity (layer)                 | 3     | 37,233         | 0.9655 $\pm$ 0.0009 | 0.020 $\pm$ 0.001 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: abs(w) $\times$ activity (layer) | 3     | 37,233         | 0.9805 $\pm$ 0.0004 | 0.005 $\pm$ 0.001 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: synaptic drive (layer)           | 3     | 37,233         | 0.9822 $\pm$ 0.0003 | 0.004 $\pm$ 0.000 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | prune-during: random (layer)                   | 3     | 37,233         | 0.9339 $\pm$ 0.0015 | 0.034 $\pm$ 0.002 | $7.45 \times 10^4$ | $2.73 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.02    | SET (dynamic, random regrow)                   | 3     | 37,233         | 0.9743 $\pm$ 0.0006 | 0.005 $\pm$ 0.000 | $7.45 \times 10^4$ | $2.01 \times 10^{11}$ | 467683 $\pm$ 20709 |
| 0.02    | RigL (dynamic, grad regrow)                    | 3     | 37,233         | 0.9785 $\pm$ 0.0004 | 0.005 $\pm$ 0.001 | $7.45 \times 10^4$ | $2.01 \times 10^{11}$ | 245749 $\pm$ 12815 |
| 0.02    | static random sparse                           | 3     | 37,233         | 0.9590 $\pm$ 0.0013 | 0.009 $\pm$ 0.001 | $7.45 \times 10^4$ | $2.01 \times 10^{11}$ | -                  |
| 0.01    | dense small (additive)                         | 3     | 18,791         | 0.9630 $\pm$ 0.0012 | 0.004 $\pm$ 0.001 | $3.76 \times 10^4$ | $1.01 \times 10^{11}$ | -                  |
| 0.01    | train-then-prune + finetune                    | 3     | 18,616         | 0.9752 $\pm$ 0.0020 | 0.006 $\pm$ 0.000 | $3.72 \times 10^4$ | $1.01 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: magnitude (layer)                | 3     | 18,616         | 0.9823 $\pm$ 0.0001 | 0.007 $\pm$ 0.000 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: magnitude (global)               | 3     | 18,616         | 0.9844 $\pm$ 0.0003 | 0.004 $\pm$ 0.000 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: activity (layer)                 | 3     | 18,616         | 0.9366 $\pm$ 0.0015 | 0.037 $\pm$ 0.001 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: abs(w) $\times$ activity (layer) | 3     | 18,616         | 0.9725 $\pm$ 0.0008 | 0.012 $\pm$ 0.001 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: synaptic drive (layer)           | 3     | 18,616         | 0.9756 $\pm$ 0.0010 | 0.012 $\pm$ 0.000 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | prune-during: random (layer)                   | 3     | 18,616         | 0.8913 $\pm$ 0.0060 | 0.071 $\pm$ 0.003 | $3.72 \times 10^4$ | $2.66 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.01    | SET (dynamic, random regrow)                   | 3     | 18,616         | 0.9603 $\pm$ 0.0010 | 0.008 $\pm$ 0.001 | $3.72 \times 10^4$ | $1.01 \times 10^{11}$ | -                  |
| 0.01    | RigL (dynamic, grad regrow)                    | 3     | 18,616         | 0.9741 $\pm$ 0.0016 | 0.004 $\pm$ 0.001 | $3.72 \times 10^4$ | $1.01 \times 10^{11}$ | 444658 $\pm$ 67026 |
| 0.01    | static random sparse                           | 3     | 18,616         | 0.9259 $\pm$ 0.0034 | 0.015 $\pm$ 0.001 | $3.72 \times 10^4$ | $1.01 \times 10^{11}$ | -                  |
| 0.005   | dense small (additive)                         | 3     | 9,672          | 0.9407 $\pm$ 0.0010 | 0.009 $\pm$ 0.001 | $1.93 \times 10^4$ | $5.22 \times 10^{10}$ | -                  |
| 0.005   | train-then-prune + finetune                    | 3     | 9,308          | 0.9451 $\pm$ 0.0060 | 0.017 $\pm$ 0.006 | $1.86 \times 10^4$ | $1.01 \times 10^{13}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: magnitude (layer)                | 3     | 9,308          | 0.9704 $\pm$ 0.0005 | 0.015 $\pm$ 0.001 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: magnitude (global)               | 3     | 9,308          | 0.9781 $\pm$ 0.0002 | 0.006 $\pm$ 0.000 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: activity (layer)                 | 3     | 9,322          | 0.8885 $\pm$ 0.0109 | 0.100 $\pm$ 0.037 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: abs(w) $\times$ activity (layer) | 3     | 9,308          | 0.9522 $\pm$ 0.0013 | 0.038 $\pm$ 0.016 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: synaptic drive (layer)           | 3     | 9,308          | 0.9603 $\pm$ 0.0020 | 0.025 $\pm$ 0.000 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | prune-during: random (layer)                   | 3     | 9,308          | 0.7184 $\pm$ 0.0230 | 0.157 $\pm$ 0.017 | $1.86 \times 10^4$ | $2.62 \times 10^{12}$ | 87336 $\pm$ 7781   |
| 0.005   | SET (dynamic, random regrow)                   | 3     | 9,308          | 0.9365 $\pm$ 0.0011 | 0.012 $\pm$ 0.001 | $1.86 \times 10^4$ | $5.03 \times 10^{10}$ | -                  |
| 0.005   | RigL (dynamic, grad regrow)                    | 3     | 9,308          | 0.9631 $\pm$ 0.0027 | 0.005 $\pm$ 0.001 | $1.86 \times 10^4$ | $5.03 \times 10^{10}$ | -                  |
| 0.005   | static random sparse                           | 3     | 9,308          | 0.7271 $\pm$ 0.0271 | 0.053 $\pm$ 0.012 | $1.86 \times 10^4$ | $5.03 \times 10^{10}$ | -                  |

Table 7: MNIST MLP (784-1024-1024-10, 1.86M weights), all arms and budgets, 3 seeds. Density is the final active fraction of the over-parameterised network; dense small is a width-matched network at the same budget. Training FLOPs count three forward passes per step at the actual density of that step.

| density | arm                                                       | seeds | active weights | test acc            | ECE               | infer FLOPs        | cum. train FLOPs      | samples to 85%      |
|---------|-----------------------------------------------------------|-------|----------------|---------------------|-------------------|--------------------|-----------------------|---------------------|
| 1       | dense big (no pruning)                                    | 3     | 2,195,648      | 0.9133 $\pm$ 0.0010 | 0.031 $\pm$ 0.001 | $3.08 \times 10^8$ | $9.23 \times 10^{14}$ | 411367 $\pm$ 15611  |
| 0.1     | dense small (additive)                                    | 3     | 219,735        | 0.8729 $\pm$ 0.0007 | 0.029 $\pm$ 0.001 | $3.10 \times 10^7$ | $9.31 \times 10^{13}$ | 632089 $\pm$ 19636  |
| 0.1     | train-then-prune + finetune                               | 3     | 219,565        | 0.9072 $\pm$ 0.0014 | 0.036 $\pm$ 0.001 | $8.19 \times 10^7$ | $1.05 \times 10^{15}$ | 411367 $\pm$ 15611  |
| 0.1     | prune-during: magnitude (global)                          | 3     | 219,565        | 0.9028 $\pm$ 0.0029 | 0.033 $\pm$ 0.001 | $7.37 \times 10^7$ | $4.19 \times 10^{14}$ | 387708 $\pm$ 26819  |
| 0.1     | prune-during: magnitude (ERK layer budget)                | 3     | 219,565        | 0.8930 $\pm$ 0.0004 | 0.030 $\pm$ 0.001 | $2.69 \times 10^7$ | $2.80 \times 10^{14}$ | 374564 $\pm$ 22672  |
| 0.1     | prune-during: synaptic drive, per-neuron normalized (ERK) | 3     | 219,565        | 0.8759 $\pm$ 0.0033 | 0.026 $\pm$ 0.003 | $2.69 \times 10^7$ | $2.80 \times 10^{14}$ | 633096 $\pm$ 45722  |
| 0.1     | RigL (dynamic, grad regrow)                               | 3     | 219,565        | 0.8759 $\pm$ 0.0005 | 0.026 $\pm$ 0.002 | $2.69 \times 10^7$ | $8.07 \times 10^{13}$ | 625325 $\pm$ 28326  |
| 0.03    | dense small (additive)                                    | 3     | 67,905         | 0.8284 $\pm$ 0.0010 | 0.024 $\pm$ 0.001 | $9.72 \times 10^6$ | $2.92 \times 10^{13}$ | -                   |
| 0.03    | train-then-prune + finetune                               | 3     | 65,869         | 0.8838 $\pm$ 0.0022 | 0.026 $\pm$ 0.001 | $3.58 \times 10^7$ | $9.77 \times 10^{14}$ | 411367 $\pm$ 15611  |
| 0.03    | prune-during: magnitude (global)                          | 3     | 65,869         | 0.8908 $\pm$ 0.0017 | 0.028 $\pm$ 0.002 | $2.98 \times 10^7$ | $3.40 \times 10^{14}$ | 391024 $\pm$ 6160   |
| 0.03    | prune-during: magnitude (ERK layer budget)                | 3     | 65,870         | 0.8692 $\pm$ 0.0026 | 0.023 $\pm$ 0.003 | $8.38 \times 10^6$ | $2.39 \times 10^{14}$ | 495322 $\pm$ 145564 |
| 0.03    | prune-during: synaptic drive, per-neuron normalized (ERK) | 3     | 65,870         | 0.7928 $\pm$ 0.0048 | 0.020 $\pm$ 0.003 | $8.38 \times 10^6$ | $2.39 \times 10^{14}$ | -                   |
| 0.03    | RigL (dynamic, grad regrow)                               | 3     | 65,870         | 0.8353 $\pm$ 0.0019 | 0.018 $\pm$ 0.001 | $8.38 \times 10^6$ | $2.52 \times 10^{13}$ | -                   |
| 0.01    | dense small (additive)                                    | 3     | 22,911         | 0.7709 $\pm$ 0.0036 | 0.020 $\pm$ 0.004 | $3.32 \times 10^6$ | $9.97 \times 10^{12}$ | -                   |
| 0.01    | train-then-prune + finetune                               | 3     | 21,956         | 0.8051 $\pm$ 0.0010 | 0.015 $\pm$ 0.003 | $1.60 \times 10^7$ | $9.47 \times 10^{14}$ | 411367 $\pm$ 15611  |
| 0.01    | prune-during: magnitude (global)                          | 3     | 21,956         | 0.8570 $\pm$ 0.0020 | 0.023 $\pm$ 0.001 | $1.16 \times 10^7$ | $3.09 \times 10^{14}$ | 390029 $\pm$ 6850   |
| 0.01    | prune-during: magnitude (ERK layer budget)                | 3     | 21,957         | 0.8134 $\pm$ 0.0007 | 0.015 $\pm$ 0.001 | $2.78 \times 10^6$ | $2.27 \times 10^{14}$ | -                   |
| 0.01    | prune-during: synaptic drive, per-neuron normalized (ERK) | 3     | 21,957         | 0.6004 $\pm$ 0.0141 | 0.013 $\pm$ 0.006 | $2.78 \times 10^6$ | $2.27 \times 10^{14}$ | -                   |
| 0.01    | RigL (dynamic, grad regrow)                               | 3     | 21,957         | 0.5392 $\pm$ 0.3110 | 0.008 $\pm$ 0.006 | $2.78 \times 10^6$ | $8.35 \times 10^{12}$ | -                   |

Table 8: CIFAR-10 CNN (2.2M weights), all arms and budgets, 3 seeds. Columns as in Table 7; the data-efficiency column is the number of training samples until 85% test accuracy.

| density | arm                                        | seeds | active weights | test acc            | ECE               | infer FLOPs        | cum. train FLOPs      | samples to 85%     |
|---------|--------------------------------------------|-------|----------------|---------------------|-------------------|--------------------|-----------------------|--------------------|
| 1       | dense big (no pruning)                     | 3     | 11,164,352     | 0.9248 $\pm$ 0.0026 | 0.035 $\pm$ 0.001 | $1.11 \times 10^9$ | $3.33 \times 10^{15}$ | 436855 $\pm$ 29455 |
| 0.2     | dense small (additive)                     | 2     | 2,239,484      | 0.9093 $\pm$ 0.0000 | 0.035 $\pm$ 0.002 | $2.25 \times 10^8$ | $6.74 \times 10^{14}$ | 514366 $\pm$ 32163 |
| 0.2     | prune-during: magnitude (layer)            | 2     | 2,232,871      | 0.9179 $\pm$ 0.0009 | 0.038 $\pm$ 0.000 | $2.22 \times 10^8$ | $1.35 \times 10^{15}$ | 382401 $\pm$ 5235  |
| 0.2     | prune-during: magnitude (global)           | 2     | 2,232,870      | 0.9189 $\pm$ 0.0013 | 0.038 $\pm$ 0.000 | $3.45 \times 10^8$ | $1.65 \times 10^{15}$ | 377044 $\pm$ 750   |
| 0.05    | dense small (additive)                     | 3     | 562,229        | 0.8866 $\pm$ 0.0018 | 0.030 $\pm$ 0.002 | $5.58 \times 10^7$ | $1.67 \times 10^{14}$ | 665417 $\pm$ 15510 |
| 0.05    | train-then-prune + finetune                | 3     | 558,218        | 0.9227 $\pm$ 0.0006 | 0.035 $\pm$ 0.001 | $2.71 \times 10^8$ | $3.74 \times 10^{15}$ | 434143 $\pm$ 30338 |
| 0.05    | prune-during: magnitude (layer)            | 2     | 558,217        | 0.9108 $\pm$ 0.0002 | 0.037 $\pm$ 0.001 | $5.55 \times 10^7$ | $9.79 \times 10^{14}$ | 387334 $\pm$ 8668  |
| 0.05    | prune-during: magnitude (global)           | 3     | 558,218        | 0.9196 $\pm$ 0.0023 | 0.034 $\pm$ 0.001 | $1.85 \times 10^8$ | $1.33 \times 10^{15}$ | 384521 $\pm$ 15853 |
| 0.05    | prune-during: magnitude (ERK layer budget) | 3     | 558,213        | 0.9203 $\pm$ 0.0009 | 0.034 $\pm$ 0.001 | $1.39 \times 10^8$ | $1.43 \times 10^{15}$ | 399622 $\pm$ 18599 |
| 0.05    | RigL (dynamic, grad regrow)                | 3     | 558,213        | 0.9097 $\pm$ 0.0016 | 0.034 $\pm$ 0.001 | $1.39 \times 10^8$ | $4.17 \times 10^{14}$ | 524129 $\pm$ 20642 |
| 0.02    | dense small (additive)                     | 3     | 225,892        | 0.8647 $\pm$ 0.0030 | 0.026 $\pm$ 0.000 | $2.25 \times 10^7$ | $6.76 \times 10^{13}$ | 801128 $\pm$ 18275 |
| 0.02    | train-then-prune + finetune                | 3     | 223,287        | 0.9094 $\pm$ 0.0018 | 0.027 $\pm$ 0.002 | $1.49 \times 10^8$ | $3.56 \times 10^{15}$ | 434143 $\pm$ 30338 |
| 0.02    | prune-during: magnitude (global)           | 3     | 223,287        | 0.9146 $\pm$ 0.0015 | 0.034 $\pm$ 0.002 | $1.07 \times 10^8$ | $1.20 \times 10^{15}$ | 388422 $\pm$ 24257 |
| 0.02    | prune-during: magnitude (ERK layer budget) | 3     | 223,288        | 0.9125 $\pm$ 0.0012 | 0.032 $\pm$ 0.001 | $5.71 \times 10^7$ | $1.26 \times 10^{15}$ | 394149 $\pm$ 11727 |
| 0.02    | RigL (dynamic, grad regrow)                | 3     | 223,288        | 0.8950 $\pm$ 0.0018 | 0.030 $\pm$ 0.003 | $5.71 \times 10^7$ | $1.71 \times 10^{14}$ | 597590 $\pm$ 17418 |
| 0.005   | dense small (additive)                     | 3     | 55,872         | 0.8131 $\pm$ 0.0026 | 0.021 $\pm$ 0.001 | $6.12 \times 10^6$ | $1.84 \times 10^{13}$ | -                  |
| 0.005   | train-then-prune + finetune                | 3     | 55,822         | 0.8364 $\pm$ 0.0015 | 0.016 $\pm$ 0.001 | $5.27 \times 10^7$ | $3.41 \times 10^{15}$ | 445811 $\pm$ 28637 |
| 0.005   | prune-during: magnitude (layer)            | 2     | 55,819         | 0.8311 $\pm$ 0.0006 | 0.022 $\pm$ 0.001 | $5.55 \times 10^6$ | $8.67 \times 10^{14}$ | 392455 $\pm$ 2533  |
| 0.005   | prune-during: magnitude (global)           | 3     | 55,822         | 0.8896 $\pm$ 0.0003 | 0.027 $\pm$ 0.001 | $3.40 \times 10^7$ | $1.09 \times 10^{15}$ | 401879 $\pm$ 17074 |
| 0.005   | prune-during: magnitude (ERK layer budget) | 3     | 55,820         | 0.8783 $\pm$ 0.0022 | 0.024 $\pm$ 0.001 | $1.38 \times 10^7$ | $1.18 \times 10^{15}$ | 391071 $\pm$ 8159  |
| 0.005   | RigL (dynamic, grad regrow)                | 3     | 55,820         | 0.8486 $\pm$ 0.0030 | 0.026 $\pm$ 0.001 | $1.38 \times 10^7$ | $4.15 \times 10^{13}$ | 907936             |

Table 9: CIFAR-10 ResNet-18 (11.2M weights), all arms and budgets. Columns as in Table 7; samples to 85% test accuracy.

| density | arm                                                     | seeds | active conv weights | init acc            | final acc           | best acc            | ECE               | infer FLOPs        | adaptation train FLOPs | samples to 90%     |
|---------|---------------------------------------------------------|-------|---------------------|---------------------|---------------------|---------------------|-------------------|--------------------|------------------------|--------------------|
| 1       | pre-trained, dense fine-tune (no pruning)               | 2     | 11,166,912          | 0.1085 $\pm$ 0.0170 | 0.9593 $\pm$ 0.0000 | 0.9598 $\pm$ 0.0000 | 0.018 $\pm$ 0.001 | 1.18 $\times 10^9$ | 1.78 $\times 10^{15}$  | 52515 $\pm$ 1566   |
| 0.05    | from scratch, dense small (width-scaled)                | 2     | 562,759             | 0.1000 $\pm$ 0.0000 | 0.8484 $\pm$ 0.0004 | 0.8486 $\pm$ 0.0006 | 0.022 $\pm$ 0.001 | 7.19 $\times 10^7$ | 1.08 $\times 10^{14}$  | -                  |
| 0.05    | from scratch, prune-during (global magnitude)           | 2     | 558,346             | 0.1000 $\pm$ 0.0000 | 0.8835 $\pm$ 0.0007 | 0.8838 $\pm$ 0.0010 | 0.031 $\pm$ 0.001 | 2.40 $\times 10^8$ | 7.72 $\times 10^{14}$  | -                  |
| 0.05    | pre-trained, one-shot prune then fine-tune              | 2     | 558,346             | 0.1000 $\pm$ 0.0000 | 0.9306 $\pm$ 0.0012 | 0.9316 $\pm$ 0.0007 | 0.020 $\pm$ 0.001 | 2.01 $\times 10^8$ | 3.02 $\times 10^{14}$  | 165526 $\pm$ 4694  |
| 0.05    | pre-trained, random ERK mask + RgL regrow               | 2     | 558,347             | 0.1000 $\pm$ 0.0000 | 0.8690 $\pm$ 0.0021 | 0.8692 $\pm$ 0.0020 | 0.023 $\pm$ 0.001 | 1.78 $\times 10^8$ | 2.68 $\times 10^{14}$  | -                  |
| 0.05    | pre-trained, prune-during adaptation (ERK layer budget) | 2     | 558,347             | 0.1085 $\pm$ 0.0170 | 0.9343 $\pm$ 0.0005 | 0.9356 $\pm$ 0.0003 | 0.015 $\pm$ 0.001 | 1.78 $\times 10^8$ | 8.10 $\times 10^{14}$  | 52380 $\pm$ 1430   |
| 0.05    | pre-trained, prune-during adaptation (global magnitude) | 2     | 558,346             | 0.1085 $\pm$ 0.0170 | 0.9365 $\pm$ 0.0008 | 0.9377 $\pm$ 0.0005 | 0.017 $\pm$ 0.001 | 2.03 $\times 10^8$ | 7.26 $\times 10^{14}$  | 52276 $\pm$ 1327   |
| 0.02    | from scratch, dense small (width-scaled)                | 2     | 226,242             | 0.0991 $\pm$ 0.0048 | 0.8240 $\pm$ 0.0006 | 0.8241 $\pm$ 0.0005 | 0.022 $\pm$ 0.001 | 3.29 $\times 10^7$ | 4.93 $\times 10^{13}$  | -                  |
| 0.02    | from scratch, prune-during (global magnitude)           | 2     | 223,338             | 0.1000 $\pm$ 0.0000 | 0.8747 $\pm$ 0.0016 | 0.8752 $\pm$ 0.0011 | 0.031 $\pm$ 0.001 | 1.44 $\times 10^8$ | 6.94 $\times 10^{14}$  | -                  |
| 0.02    | pre-trained, one-shot prune then fine-tune              | 2     | 223,338             | 0.1000 $\pm$ 0.0000 | 0.9029 $\pm$ 0.0010 | 0.9038 $\pm$ 0.0000 | 0.023 $\pm$ 0.000 | 1.23 $\times 10^8$ | 1.85 $\times 10^{14}$  | 401984 $\pm$ 12502 |
| 0.02    | pre-trained, random ERK mask + RgL regrow               | 2     | 223,341             | 0.1000 $\pm$ 0.0000 | 0.8355 $\pm$ 0.0020 | 0.8355 $\pm$ 0.0019 | 0.025 $\pm$ 0.000 | 7.17 $\times 10^7$ | 1.08 $\times 10^{14}$  | -                  |
| 0.02    | pre-trained, prune-during adaptation (ERK layer budget) | 2     | 223,341             | 0.1085 $\pm$ 0.0170 | 0.8966 $\pm$ 0.0019 | 0.9292 $\pm$ 0.0008 | 0.020 $\pm$ 0.000 | 7.17 $\times 10^7$ | 7.10 $\times 10^{14}$  | 52567 $\pm$ 1617   |
| 0.005   | pre-trained, prune-during adaptation (global magnitude) | 2     | 223,338             | 0.1085 $\pm$ 0.0170 | 0.9098 $\pm$ 0.0008 | 0.9329 $\pm$ 0.0010 | 0.019 $\pm$ 0.001 | 1.25 $\times 10^8$ | 6.57 $\times 10^{14}$  | 52360 $\pm$ 1410   |
| 0.005   | from scratch, dense small (width-scaled)                | 2     | 56,112              | 0.1000 $\pm$ 0.0000 | 0.7654 $\pm$ 0.0001 | 0.7664 $\pm$ 0.0011 | 0.022 $\pm$ 0.001 | 1.19 $\times 10^7$ | 1.78 $\times 10^{13}$  | -                  |
| 0.005   | from scratch, prune-during (global magnitude)           | 2     | 55,835              | 0.1000 $\pm$ 0.0000 | 0.8380 $\pm$ 0.0037 | 0.8384 $\pm$ 0.0034 | 0.028 $\pm$ 0.002 | 5.39 $\times 10^7$ | 6.27 $\times 10^{14}$  | -                  |
| 0.005   | pre-trained, one-shot prune then fine-tune              | 2     | 55,835              | 0.0884 $\pm$ 0.0116 | 0.8171 $\pm$ 0.0014 | 0.8184 $\pm$ 0.0019 | 0.019 $\pm$ 0.002 | 5.64 $\times 10^7$ | 8.45 $\times 10^{13}$  | -                  |
| 0.005   | pre-trained, random ERK mask + RgL regrow               | 2     | 55,834              | 0.1000 $\pm$ 0.0000 | 0.7682 $\pm$ 0.0037 | 0.7690 $\pm$ 0.0044 | 0.023 $\pm$ 0.002 | 1.79 $\times 10^7$ | 2.69 $\times 10^{13}$  | -                  |
| 0.005   | pre-trained, prune-during adaptation (ERK layer budget) | 2     | 55,834              | 0.1085 $\pm$ 0.0170 | 0.7883 $\pm$ 0.0022 | 0.9314 $\pm$ 0.0006 | 0.020 $\pm$ 0.000 | 1.79 $\times 10^7$ | 6.59 $\times 10^{14}$  | 52422 $\pm$ 1473   |
| 0.005   | pre-trained, prune-during adaptation (global magnitude) | 2     | 55,835              | 0.1085 $\pm$ 0.0170 | 0.8087 $\pm$ 0.0012 | 0.9306 $\pm$ 0.0002 | 0.016 $\pm$ 0.000 | 5.68 $\times 10^7$ | 6.04 $\times 10^{14}$  | 52422 $\pm$ 1473   |

Table 10: ImageNet-pre-trained ResNet-18 adapted to CIFAR-10 (128 $\times$ 128 input, 10 epochs) while being pruned. Density is relative to the 11.17M convolutional weights; the classifier is kept dense in every arm. Init acc is the accuracy before adaptation (fresh classifier). Adaptation FLOPs exclude the ImageNet pre-training, which is common to all pre-trained arms.

| method                                                                  | seeds | avg acc (final) | forgetting      | retention     | ECE (all classes) | train FLOPs           | active params (predictor) | unknown-class AUROC (softmax) | unknown-class AUROC (fast/slow agreement) |
|-------------------------------------------------------------------------|-------|-----------------|-----------------|---------------|-------------------|-----------------------|---------------------------|-------------------------------|-------------------------------------------|
| fine-tune (no protection)                                               | 3     | 0.1969 ± 0.0004 | 0.9950 ± 0.0001 | 0.000 ± 0.000 | 0.731 ± 0.005     | $2.90 \times 10^{11}$ | 269,322                   | 0.649 ± 0.022                 | -                                         |
| experience replay (buffer 500)                                          | 3     | 0.8515 ± 0.0043 | 0.1747 ± 0.0062 | 0.824 ± 0.006 | 0.087 ± 0.006     | $5.82 \times 10^{11}$ | 269,322                   | 0.845 ± 0.007                 | -                                         |
| ER, class-balanced batches (buffer 500)                                 | 3     | 0.8431 ± 0.0004 | 0.1860 ± 0.0132 | 0.811 ± 0.013 | 0.105 ± 0.008     | $1.16 \times 10^{12}$ | 269,322                   | 0.842 ± 0.021                 | -                                         |
| CLS + sleep (replay + distill)                                          | 3     | 0.8784 ± 0.0068 | 0.1302 ± 0.0087 | 0.868 ± 0.009 | 0.069 ± 0.005     | $2.87 \times 10^{11}$ | 269,322                   | 0.855 ± 0.017                 | -                                         |
| CLS + sleep (no distill)                                                | 3     | 0.8497 ± 0.0059 | 0.2542 ± 0.0213 | 0.744 ± 0.021 | 0.123 ± 0.028     | $1.84 \times 10^{12}$ | 269,322                   | 0.774 ± 0.015                 | 0.511 ± 0.012                             |
| CLS + sleep + downscale 0.1                                             | 3     | 0.7744 ± 0.0211 | 0.2700 ± 0.0273 | 0.728 ± 0.028 | 0.090 ± 0.003     | $1.84 \times 10^{12}$ | 269,322                   | 0.834 ± 0.004                 | 0.508 ± 0.002                             |
| CLS + sleep + prune 5%/sleep (magnitude)                                | 3     | 0.7920 ± 0.0113 | 0.2485 ± 0.0158 | 0.750 ± 0.016 | 0.133 ± 0.027     | $1.84 \times 10^{12}$ | 269,322                   | 0.776 ± 0.013                 | 0.512 ± 0.012                             |
| CLS + sleep + prune 5%/sleep (activity)                                 | 3     | 0.7809 ± 0.0156 | 0.2582 ± 0.0191 | 0.739 ± 0.020 | 0.118 ± 0.025     | $1.84 \times 10^{12}$ | 208,514                   | 0.776 ± 0.017                 | 0.511 ± 0.012                             |
| CLS + sleep + downscale + prune                                         | 3     | 0.7875 ± 0.0198 | 0.2540 ± 0.0247 | 0.744 ± 0.025 | 0.107 ± 0.022     | $1.84 \times 10^{12}$ | 208,514                   | 0.744 ± 0.016                 | 0.521 ± 0.012                             |
| CLS2: balanced sleep replay, no distill, dense fast [predict: fast]     | 3     | 0.1969 ± 0.0002 | 0.9964 ± 0.0004 | 0.000 ± 0.000 | 0.121 ± 0.024     | $1.84 \times 10^{12}$ | 208,514                   | 0.771 ± 0.022                 | 0.510 ± 0.013                             |
| CLS2: balanced sleep replay, no distill, dense fast [predict: max_conf] | 3     | 0.6794 ± 0.0256 | 0.3947 ± 0.0325 | 0.604 ± 0.032 | -                 | -                     | -                         | -                             | -                                         |
| CLS2: balanced sleep replay, no distill, dense fast [predict: fast]     | 3     | 0.8703 ± 0.0059 | 0.1528 ± 0.0075 | 0.846 ± 0.008 | 0.080 ± 0.005     | $3.27 \times 10^{12}$ | 269,322                   | 0.857 ± 0.017                 | 0.505 ± 0.001                             |
| CLS2 + sparse fast (k-WTA 10%) [predict: fast]                          | 3     | 0.1969 ± 0.0002 | 0.9973 ± 0.0004 | 0.000 ± 0.000 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 10%) [predict: max_conf]                      | 3     | 0.6820 ± 0.0116 | 0.3924 ± 0.0144 | 0.607 ± 0.015 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 5%) [predict: fast]                           | 3     | 0.8703 ± 0.0059 | 0.1528 ± 0.0075 | 0.846 ± 0.008 | 0.080 ± 0.005     | $3.27 \times 10^{12}$ | 269,322                   | 0.857 ± 0.017                 | 0.511 ± 0.002                             |
| CLS2 + sparse fast (k-WTA 5%) [predict: max_conf]                       | 3     | 0.1969 ± 0.0006 | 0.9965 ± 0.0009 | 0.000 ± 0.000 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 5%) [predict: max_conf]                       | 3     | 0.6595 ± 0.0190 | 0.4206 ± 0.0230 | 0.578 ± 0.023 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 5%) [predict: max_conf]                       | 3     | 0.8703 ± 0.0059 | 0.1528 ± 0.0075 | 0.846 ± 0.008 | 0.080 ± 0.005     | $3.27 \times 10^{12}$ | 269,322                   | 0.857 ± 0.017                 | 0.512 ± 0.004                             |
| CLS2 + sparse fast 10% + magnitude prune 5%/sleep [predict: fast]       | 3     | 0.1969 ± 0.0002 | 0.9973 ± 0.0004 | 0.000 ± 0.000 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast 10% + magnitude prune 5%/sleep [predict: max_conf]   | 3     | 0.6728 ± 0.0133 | 0.4030 ± 0.0168 | 0.595 ± 0.017 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast 10% + magnitude prune 5%/sleep                       | 3     | 0.8689 ± 0.0030 | 0.1548 ± 0.0039 | 0.844 ± 0.004 | 0.079 ± 0.003     | $3.27 \times 10^{12}$ | -                         | 0.857 ± 0.013                 | 0.510 ± 0.001                             |
| CLS2 + sparse fast 10% + magnitude prune 5%/sleep                       | 3     | 0.1968 ± 0.0004 | 0.9957 ± 0.0021 | 0.000 ± 0.000 | -                 | -                     | 208,514                   | -                             | -                                         |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3 [predict: fast]         | 3     | 0.6570 ± 0.0569 | 0.4214 ± 0.0716 | 0.577 ± 0.072 | -                 | -                     | -                         | -                             | -                                         |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3 [predict: max_conf]     | 3     | 0.8679 ± 0.0105 | 0.1510 ± 0.0145 | 0.848 ± 0.015 | 0.080 ± 0.007     | $2.25 \times 10^{12}$ | 269,322                   | 0.857 ± 0.006                 | 0.512 ± 0.006                             |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3                         | 3     | 0.9762 ± 0.0009 | 0.0086 ± 0.0031 | 0.991 ± 0.003 | 0.005 ± 0.002     | $8.79 \times 10^{11}$ | 269,322                   | 0.934 ± 0.005                 | -                                         |
| joint retraining (upper bound)                                          | 3     | -               | -               | -             | -                 | -                     | -                         | -                             | -                                         |

Table 11: Split MNIST, class-incremental (single 10-way head, 5 tasks of 2 classes), 3 seeds. Rows marked [predict: ...] use an alternative read-out of the same run. Unknown-class AUROC: how well the confidence separates test samples of classes not yet learned (0.5 = chance).

| method                                                                  | seeds | avg acc. (final) | forgetting      | retention     | ECE (all classes) | train FLOPs             | active params | unknown-class AUROC (softmax) | unknown-class AUROC (fast/slow agreement) |
|-------------------------------------------------------------------------|-------|------------------|-----------------|---------------|-------------------|-------------------------|---------------|-------------------------------|-------------------------------------------|
| fine-tune (no protection)                                               | 3     | 0.4894 ± 0.0151  | 0.5384 ± 0.0166 | 0.447 ± 0.017 | 0.502 ± 0.043     | 2.90 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
| experience replay (buffer 500)                                          | 3     | 0.8739 ± 0.0103  | 0.1121 ± 0.0115 | 0.885 ± 0.012 | 0.187 ± 0.042     | 5.81 × 10 <sup>14</sup> | 269,322       | -                             | -                                         |
| ER, class-balanced batches (buffer 500)                                 | 3     | 0.8271 ± 0.0142  | 0.1646 ± 0.0157 | 0.831 ± 0.016 | 0.214 ± 0.017     | 1.16 × 10 <sup>13</sup> | 269,322       | -                             | -                                         |
| CLS + sleep (replay + distill)                                          | 3     | 0.8828 ± 0.0037  | 0.0764 ± 0.0050 | 0.921 ± 0.005 | 0.140 ± 0.012     | 2.36 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
| CLS + sleep (no distill)                                                | 3     | 0.8780 ± 0.0038  | 0.0939 ± 0.0047 | 0.902 ± 0.005 | 0.110 ± 0.020     | 6.00 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
| CLS + sleep + downscale 0.1                                             | 3     | 0.8856 ± 0.0006  | 0.0847 ± 0.0002 | 0.912 ± 0.000 | 0.125 ± 0.005     | 6.00 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
| CLS + sleep + prune 5%/sleep (magnitude)                                | 3     | 0.8269 ± 0.0064  | 0.1512 ± 0.0070 | 0.843 ± 0.007 | 0.151 ± 0.022     | 6.00 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
| CLS + sleep + prune 5%/sleep (activity)                                 | 3     | 0.8731 ± 0.0040  | 0.0973 ± 0.0042 | 0.899 ± 0.004 | 0.130 ± 0.022     | 6.00 × 10 <sup>12</sup> | 161,462       | -                             | -                                         |
| CLS + sleep + downscale + prune                                         | 3     | 0.8448 ± 0.0026  | 0.1599 ± 0.0029 | 0.833 ± 0.003 | 0.195 ± 0.029     | 6.00 × 10 <sup>12</sup> | 161,462       | -                             | -                                         |
| CLS2: balanced sleep replay, no distill, dense fast [predict: fast]     | 3     | 0.8266 ± 0.0076  | 0.1382 ± 0.0088 | 0.856 ± 0.009 | 0.120 ± 0.029     | 6.00 × 10 <sup>12</sup> | 161,462       | -                             | -                                         |
| CLS2: balanced sleep replay, no distill, dense fast [predict: max_conf] | 3     | 0.4026 ± 0.0574  | 0.6295 ± 0.0630 | 0.351 ± 0.065 | -                 | -                       | -             | -                             | -                                         |
| CLS2: balanced sleep replay, no distill, dense fast [predict: fast]     | 3     | 0.8577 ± 0.0064  | 0.1256 ± 0.0070 | 0.871 ± 0.007 | -                 | -                       | 269,322       | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 10%) [predict: fast]                          | 3     | 0.8903 ± 0.0020  | 0.0662 ± 0.0024 | 0.931 ± 0.003 | 0.109 ± 0.022     | 1.18 × 10 <sup>13</sup> | -             | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 10%) [predict: max_conf]                      | 3     | 0.4661 ± 0.0150  | 0.5601 ± 0.0169 | 0.423 ± 0.017 | -                 | -                       | -             | -                             | -                                         |
| CLS2 + sparse fast (k-WTA 10%)                                          | 3     | 0.8708 ± 0.0145  | 0.1119 ± 0.0161 | 0.885 ± 0.017 | -                 | -                       | 269,322       | -                             | -                                         |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3 [predict: fast]         | 3     | 0.8903 ± 0.0020  | 0.0662 ± 0.0024 | 0.931 ± 0.003 | 0.109 ± 0.022     | 1.18 × 10 <sup>13</sup> | -             | -                             | -                                         |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3 [predict: max_conf]     | 3     | 0.4080 ± 0.0138  | 0.6248 ± 0.0162 | 0.357 ± 0.016 | -                 | -                       | -             | -                             | -                                         |
| CLS2 + sparse fast 10%, sleep 300 steps lr 1e-3                         | 3     | 0.8503 ± 0.0029  | 0.1330 ± 0.0034 | 0.863 ± 0.003 | -                 | -                       | -             | -                             | -                                         |
| joint retraining (upper bound)                                          | 3     | 0.8851 ± 0.0027  | 0.0617 ± 0.0043 | 0.935 ± 0.004 | 0.097 ± 0.022     | 9.97 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |
|                                                                         | 3     | 0.9633 ± 0.0061  | 0.0034 ± 0.0005 | 0.998 ± 0.001 | 0.007 ± 0.002     | 8.87 × 10 <sup>12</sup> | 269,322       | -                             | -                                         |

Table 12: Permuted MNIST, domain-incremental (10 tasks), 3 seeds. Columns as in Table 11.

| score     | substitute         | skip frac | test acc            | ECE               | FLOPs ratio |
|-----------|--------------------|-----------|---------------------|-------------------|-------------|
| baseline  | -                  | 0         | $0.9827 \pm 0.0003$ | $0.005 \pm 0.001$ | 1.000       |
| predicted | identity           | 0.1       | $0.9745 \pm 0.0026$ | $0.007 \pm 0.001$ | 0.935       |
| predicted | identity           | 0.2       | $0.9623 \pm 0.0073$ | $0.010 \pm 0.003$ | 0.851       |
| predicted | identity           | 0.3       | $0.9494 \pm 0.0112$ | $0.013 \pm 0.005$ | 0.767       |
| predicted | identity           | 0.4       | $0.9395 \pm 0.0149$ | $0.015 \pm 0.006$ | 0.683       |
| predicted | identity           | 0.5       | $0.9290 \pm 0.0203$ | $0.019 \pm 0.008$ | 0.616       |
| predicted | identity           | 0.6       | $0.9157 \pm 0.0247$ | $0.024 \pm 0.010$ | 0.531       |
| predicted | identity           | 0.7       | $0.9066 \pm 0.0252$ | $0.027 \pm 0.010$ | 0.447       |
| predicted | identity           | 0.8       | $0.8927 \pm 0.0177$ | $0.032 \pm 0.007$ | 0.363       |
| predicted | identity_late_only | 0.3       | $0.9807 \pm 0.0005$ | $0.005 \pm 0.000$ | 0.884       |
| predicted | identity_late_only | 0.5       | $0.9800 \pm 0.0011$ | $0.006 \pm 0.000$ | 0.808       |
| predicted | identity_late_only | 0.7       | $0.9798 \pm 0.0011$ | $0.006 \pm 0.001$ | 0.724       |
| predicted | predicted          | 0.1       | $0.9789 \pm 0.0014$ | $0.006 \pm 0.001$ | 0.935       |
| predicted | predicted          | 0.2       | $0.9730 \pm 0.0031$ | $0.006 \pm 0.000$ | 0.851       |
| predicted | predicted          | 0.3       | $0.9678 \pm 0.0057$ | $0.007 \pm 0.001$ | 0.767       |
| predicted | predicted          | 0.4       | $0.9645 \pm 0.0079$ | $0.008 \pm 0.002$ | 0.683       |
| predicted | predicted          | 0.5       | $0.9601 \pm 0.0088$ | $0.010 \pm 0.003$ | 0.616       |
| predicted | predicted          | 0.6       | $0.9557 \pm 0.0094$ | $0.012 \pm 0.003$ | 0.531       |
| predicted | predicted          | 0.7       | $0.9502 \pm 0.0079$ | $0.014 \pm 0.003$ | 0.447       |
| predicted | predicted          | 0.8       | $0.9439 \pm 0.0063$ | $0.016 \pm 0.003$ | 0.363       |
| random    | identity           | 0.1       | $0.9796 \pm 0.0006$ | $0.005 \pm 0.000$ | 0.935       |
| random    | identity           | 0.2       | $0.9747 \pm 0.0013$ | $0.007 \pm 0.002$ | 0.851       |
| random    | identity           | 0.3       | $0.9671 \pm 0.0016$ | $0.008 \pm 0.001$ | 0.767       |
| random    | identity           | 0.4       | $0.9574 \pm 0.0030$ | $0.011 \pm 0.001$ | 0.683       |
| random    | identity           | 0.5       | $0.9448 \pm 0.0016$ | $0.013 \pm 0.003$ | 0.616       |
| random    | identity           | 0.6       | $0.9266 \pm 0.0065$ | $0.020 \pm 0.003$ | 0.531       |
| random    | identity           | 0.7       | $0.9028 \pm 0.0101$ | $0.030 \pm 0.004$ | 0.447       |
| random    | identity           | 0.8       | $0.8785 \pm 0.0215$ | $0.037 \pm 0.012$ | 0.363       |
| oracle    | identity           | 0.1       | $0.9774 \pm 0.0011$ | $0.006 \pm 0.001$ | 0.935       |
| oracle    | identity           | 0.2       | $0.9679 \pm 0.0018$ | $0.008 \pm 0.001$ | 0.851       |
| oracle    | identity           | 0.3       | $0.9557 \pm 0.0027$ | $0.011 \pm 0.003$ | 0.767       |
| oracle    | identity           | 0.4       | $0.9446 \pm 0.0050$ | $0.015 \pm 0.003$ | 0.683       |
| oracle    | identity           | 0.5       | $0.9353 \pm 0.0039$ | $0.018 \pm 0.003$ | 0.616       |
| oracle    | identity           | 0.6       | $0.9242 \pm 0.0056$ | $0.021 \pm 0.002$ | 0.531       |
| oracle    | identity           | 0.7       | $0.9121 \pm 0.0081$ | $0.025 \pm 0.004$ | 0.447       |
| oracle    | identity           | 0.8       | $0.8967 \pm 0.0126$ | $0.030 \pm 0.008$ | 0.363       |

Table 13: MNIST ViT: post-hoc token skipping without fine-tuning, 3 seeds. Score: which tokens are skipped (predicted change, random, or the oracle true residual norm); substitute: what replaces a skipped token (identity, identity in the late blocks only, or the predicted residual).

| score     | substitute           | skip frac | test acc | ECE   | FLOPs ratio |
|-----------|----------------------|-----------|----------|-------|-------------|
| baseline  | -                    | 0         | 0.8186   | 0.047 | 1.000       |
| predicted | identity             | 0.1       | 0.7851   | 0.054 | 0.936       |
| predicted | identity             | 0.2       | 0.7269   | 0.058 | 0.845       |
| predicted | identity             | 0.3       | 0.6647   | 0.095 | 0.768       |
| predicted | identity             | 0.4       | 0.5876   | 0.143 | 0.677       |
| predicted | identity             | 0.5       | 0.5149   | 0.194 | 0.600       |
| predicted | identity             | 0.6       | 0.4349   | 0.252 | 0.522       |
| predicted | identity             | 0.7       | 0.3510   | 0.317 | 0.432       |
| predicted | identity             | 0.8       | 0.3003   | 0.343 | 0.354       |
| predicted | identity__late__only | 0.3       | 0.8012   | 0.050 | 0.884       |
| predicted | identity__late__only | 0.5       | 0.7854   | 0.056 | 0.800       |
| predicted | identity__late__only | 0.7       | 0.7685   | 0.055 | 0.716       |
| predicted | predicted            | 0.1       | 0.8020   | 0.052 | 0.936       |
| predicted | predicted            | 0.2       | 0.7685   | 0.046 | 0.845       |
| predicted | predicted            | 0.3       | 0.7300   | 0.053 | 0.768       |
| predicted | predicted            | 0.4       | 0.6843   | 0.069 | 0.677       |
| predicted | predicted            | 0.5       | 0.6373   | 0.093 | 0.600       |
| predicted | predicted            | 0.6       | 0.5735   | 0.135 | 0.522       |
| predicted | predicted            | 0.7       | 0.5079   | 0.173 | 0.432       |
| predicted | predicted            | 0.8       | 0.4422   | 0.218 | 0.354       |
| random    | identity             | 0.1       | 0.7982   | 0.051 | 0.936       |
| random    | identity             | 0.2       | 0.7663   | 0.051 | 0.845       |
| random    | identity             | 0.3       | 0.7245   | 0.053 | 0.768       |
| random    | identity             | 0.4       | 0.6544   | 0.093 | 0.677       |
| random    | identity             | 0.5       | 0.5854   | 0.140 | 0.600       |
| random    | identity             | 0.6       | 0.5094   | 0.194 | 0.522       |
| random    | identity             | 0.7       | 0.4228   | 0.250 | 0.432       |
| random    | identity             | 0.8       | 0.3529   | 0.299 | 0.354       |
| oracle    | identity             | 0.1       | 0.7941   | 0.054 | 0.936       |
| oracle    | identity             | 0.2       | 0.7462   | 0.060 | 0.845       |
| oracle    | identity             | 0.3       | 0.6953   | 0.077 | 0.768       |
| oracle    | identity             | 0.4       | 0.6269   | 0.121 | 0.677       |
| oracle    | identity             | 0.5       | 0.5531   | 0.176 | 0.600       |
| oracle    | identity             | 0.6       | 0.4667   | 0.237 | 0.522       |
| oracle    | identity             | 0.7       | 0.3783   | 0.306 | 0.432       |
| oracle    | identity             | 0.8       | 0.3347   | 0.312 | 0.354       |

Table 14: CIFAR-10 ViT: post-hoc token skipping without fine-tuning, one seed. Columns as in Table 13.

| train mode / late attn                            | seeds | no skip             | warm-up only, skip $s_{max}$ | skip 0              | skip 0.3            | skip 0.5            | skip 0.7            | skip 0.8            | skip 0.9            | FLOPs ratio (by skip level)             |
|---------------------------------------------------|-------|---------------------|------------------------------|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|-----------------------------------------|
| layerwise / late attn full                        | 3     | 0.9827 $\pm$ 0.0003 | 0.9806 $\pm$ 0.0004          | 0.9785 $\pm$ 0.0012 | 0.9787 $\pm$ 0.0016 | 0.9788 $\pm$ 0.0015 | 0.9788 $\pm$ 0.0016 | 0.9787 $\pm$ 0.0016 | 0.9787 $\pm$ 0.0015 | 1.01 / 0.89 / 0.81 / 0.73 / 0.69 / 0.65 |
| (eval thalamic)                                   | 3     | —                   | —                            | —                   | 0.9788 $\pm$ 0.0015 | 0.9789 $\pm$ 0.0016 | 0.9786 $\pm$ 0.0015 | 0.9786 $\pm$ 0.0013 | 0.9784 $\pm$ 0.0013 | —                                       |
| (eval random)                                     | 3     | —                   | —                            | —                   | 0.9783 $\pm$ 0.0015 | 0.9787 $\pm$ 0.0009 | 0.9786 $\pm$ 0.0014 | 0.9784 $\pm$ 0.0014 | 0.9784 $\pm$ 0.0014 | —                                       |
| random / late attn full                           | 3     | 0.9827 $\pm$ 0.0003 | 0.9809 $\pm$ 0.0003          | 0.9794 $\pm$ 0.0020 | 0.9791 $\pm$ 0.0019 | 0.9792 $\pm$ 0.0015 | 0.9790 $\pm$ 0.0016 | 0.9787 $\pm$ 0.0016 | 0.9787 $\pm$ 0.0016 | 1.01 / 0.89 / 0.81 / 0.73 / 0.69 / 0.65 |
| (eval thalamic)                                   | 3     | —                   | —                            | —                   | 0.9795 $\pm$ 0.0019 | 0.9793 $\pm$ 0.0021 | 0.9793 $\pm$ 0.0021 | 0.9791 $\pm$ 0.0021 | 0.9790 $\pm$ 0.0017 | —                                       |
| (eval layerwise)                                  | 3     | —                   | —                            | —                   | 0.9789 $\pm$ 0.0016 | 0.9789 $\pm$ 0.0018 | 0.9790 $\pm$ 0.0018 | 0.9789 $\pm$ 0.0016 | 0.9789 $\pm$ 0.0017 | —                                       |
| thalamic / late attn full / identity substitution | 3     | 0.9827 $\pm$ 0.0003 | 0.9805 $\pm$ 0.0015          | 0.9785 $\pm$ 0.0015 | 0.9788 $\pm$ 0.0018 | 0.9789 $\pm$ 0.0017 | 0.9789 $\pm$ 0.0018 | 0.9787 $\pm$ 0.0015 | 0.9787 $\pm$ 0.0017 | 1.01 / 0.89 / 0.81 / 0.73 / 0.69 / 0.65 |
| (eval layerwise)                                  | 3     | —                   | —                            | —                   | 0.9787 $\pm$ 0.0017 | 0.9788 $\pm$ 0.0018 | 0.9788 $\pm$ 0.0014 | 0.9788 $\pm$ 0.0015 | 0.9788 $\pm$ 0.0016 | —                                       |
| (eval random)                                     | 3     | —                   | —                            | —                   | 0.9789 $\pm$ 0.0016 | 0.9790 $\pm$ 0.0020 | 0.9785 $\pm$ 0.0015 | 0.9790 $\pm$ 0.0016 | 0.9790 $\pm$ 0.0018 | —                                       |
| thalamic / late attn full                         | 3     | 0.9827 $\pm$ 0.0003 | 0.9810 $\pm$ 0.0010          | 0.9788 $\pm$ 0.0015 | 0.9790 $\pm$ 0.0019 | 0.9787 $\pm$ 0.0018 | 0.9786 $\pm$ 0.0016 | 0.9783 $\pm$ 0.0014 | 0.9783 $\pm$ 0.0012 | 1.01 / 0.89 / 0.81 / 0.73 / 0.69 / 0.65 |
| (eval layerwise)                                  | 3     | —                   | —                            | —                   | 0.9783 $\pm$ 0.0019 | 0.9782 $\pm$ 0.0021 | 0.9781 $\pm$ 0.0019 | 0.9780 $\pm$ 0.0018 | 0.9779 $\pm$ 0.0016 | —                                       |
| (eval random)                                     | 3     | —                   | —                            | —                   | 0.9784 $\pm$ 0.0018 | 0.9788 $\pm$ 0.0013 | 0.9786 $\pm$ 0.0015 | 0.9780 $\pm$ 0.0013 | 0.9781 $\pm$ 0.0012 | —                                       |
| thalamic / late attn pre                          | 3     | 0.9168 $\pm$ 0.0088 | 0.9140 $\pm$ 0.0089          | 0.9773 $\pm$ 0.0021 | 0.9774 $\pm$ 0.0027 | 0.9773 $\pm$ 0.0029 | 0.9773 $\pm$ 0.0027 | 0.9771 $\pm$ 0.0026 | 0.9773 $\pm$ 0.0022 | 0.99 / 0.87 / 0.80 / 0.72 / 0.68 / 0.64 |
| (eval layerwise)                                  | 3     | —                   | —                            | —                   | 0.9776 $\pm$ 0.0026 | 0.9773 $\pm$ 0.0025 | 0.9772 $\pm$ 0.0024 | 0.9771 $\pm$ 0.0023 | 0.9769 $\pm$ 0.0024 | —                                       |
| (eval random)                                     | 3     | —                   | —                            | —                   | 0.9771 $\pm$ 0.0019 | 0.9772 $\pm$ 0.0025 | 0.9772 $\pm$ 0.0020 | 0.9771 $\pm$ 0.0022 | 0.9772 $\pm$ 0.0019 | —                                       |

Table 15: MNIST ViT: token skipping in the late blocks. Skipped tokens receive the predicted residual, except in the rows marked identity substitution, where they are left unchanged. Each cell is test accuracy at the given skip ratio after fine-tuning that ramps the skip ratio from 0 to 0.7; the warm-up-only column is measured after the one-epoch warm-up of predictor and router with the ViT frozen, at the maximum skip ratio and before fine-tuning. Rows marked (eval thalamic) evaluate the same model with the thalamic router choosing the skipped tokens. For the pre-routed variant (late attn pre) the no-skip and warm-up columns are measured right after the late-block attention is swapped for pre-routed attention and before fine-tuning, which is why they are low; the skip columns are after fine-tuning.

| train mode / late attn                            | seeds | no skip | warm-up only, skip $s_{\max}$ | skip 0 | skip 0.3 | skip 0.5 | skip 0.7 | skip 0.8 | skip 0.9 | FLOPs ratio (by skip level)             |
|---------------------------------------------------|-------|---------|-------------------------------|--------|----------|----------|----------|----------|----------|-----------------------------------------|
| layerwise / late attn full                        | 1     | 0.8184  | 0.7795                        | 0.8011 | 0.8040   | 0.8028   | 0.8041   | 0.8008   | 0.7979   | 1.00 / 0.88 / 0.80 / 0.72 / 0.68 / 0.64 |
| (eval thalamic)                                   | 1     | –       | –                             | –      | 0.8058   | 0.8065   | 0.8049   | 0.8043   | 0.8006   | –                                       |
| (eval random)                                     | 1     | –       | –                             | –      | 0.8059   | 0.8067   | 0.8069   | 0.8042   | 0.7986   | –                                       |
| random / late attn full                           | 1     | 0.8184  | 0.7938                        | 0.8067 | 0.8117   | 0.8091   | 0.8072   | 0.8007   | 0.7977   | 1.00 / 0.88 / 0.80 / 0.72 / 0.68 / 0.64 |
| (eval thalamic)                                   | 1     | –       | –                             | –      | 0.8069   | 0.8080   | 0.8035   | 0.8011   | 0.7952   | –                                       |
| (eval layerwise)                                  | 1     | –       | –                             | –      | 0.8071   | 0.8026   | 0.8004   | 0.7965   | 0.7935   | –                                       |
| thalamic / late attn full / identity substitution | 1     | 0.8184  | 0.7764                        | 0.7913 | 0.8032   | 0.8082   | 0.8079   | 0.8035   | 0.7987   | 1.00 / 0.88 / 0.80 / 0.72 / 0.68 / 0.64 |
| (eval layerwise)                                  | 1     | –       | –                             | –      | 0.7999   | 0.8003   | 0.8004   | 0.7963   | 0.7905   | –                                       |
| (eval random)                                     | 1     | –       | –                             | –      | 0.8041   | 0.8072   | 0.8071   | 0.8031   | 0.7953   | –                                       |
| thalamic / late attn full                         | 1     | 0.8184  | 0.7860                        | 0.8018 | 0.8098   | 0.8110   | 0.8079   | 0.8053   | 0.8015   | 1.00 / 0.88 / 0.80 / 0.72 / 0.68 / 0.64 |
| (eval layerwise)                                  | 1     | –       | –                             | –      | 0.8090   | 0.8065   | 0.8039   | 0.8031   | 0.7976   | –                                       |
| (eval random)                                     | 1     | –       | –                             | –      | 0.8091   | 0.8110   | 0.8096   | 0.8040   | 0.7989   | –                                       |
| thalamic / late attn pre                          | 1     | 0.5679  | 0.5685                        | 0.8006 | 0.8039   | 0.8028   | 0.8015   | 0.8026   | 0.7942   | 0.99 / 0.87 / 0.79 / 0.71 / 0.68 / 0.63 |
| (eval layerwise)                                  | 1     | –       | –                             | –      | 0.8029   | 0.8020   | 0.7973   | 0.7961   | 0.7935   | –                                       |
| (eval random)                                     | 1     | –       | –                             | –      | 0.8053   | 0.8024   | 0.8044   | 0.8003   | 0.7939   | –                                       |

Table 16: CIFAR-10 ViT: token skipping in the late blocks, as in Table 15. late attn pre replaces the late-block attention by pre-routed attention (25% of keys).

| mode | keep ratio | k (keys/query) | seeds | test acc        | ECE           | attn FLOPs/block (dense) | attn FLOPs/block (sparsity exploited) | total FLOPs ratio |
|------|------------|----------------|-------|-----------------|---------------|--------------------------|---------------------------------------|-------------------|
| full | 1          | 50             | 3     | 0.9794 ± 0.0011 | 0.005 ± 0.001 | $1.28 \times 10^6$       | $1.28 \times 10^6$                    | 1.000             |
| post | 0.5        | 25             | 3     | 0.9798 ± 0.0010 | 0.004 ± 0.000 | $1.28 \times 10^6$       | $9.60 \times 10^5$                    | 0.985             |
| post | 0.25       | 13             | 3     | 0.9808 ± 0.0011 | 0.005 ± 0.001 | $1.28 \times 10^6$       | $8.06 \times 10^5$                    | 0.977             |
| post | 0.1        | 5              | 3     | 0.9814 ± 0.0001 | 0.004 ± 0.000 | $1.28 \times 10^6$       | $7.04 \times 10^5$                    | 0.973             |
| pre  | 0.5        | 25             | 3     | 0.9778 ± 0.0012 | 0.006 ± 0.000 | $8.00 \times 10^5$       | $8.00 \times 10^5$                    | 0.977             |
| pre  | 0.25       | 13             | 3     | 0.9804 ± 0.0007 | 0.004 ± 0.001 | $4.93 \times 10^5$       | $4.93 \times 10^5$                    | 0.962             |
| pre  | 0.1        | 5              | 3     | 0.9808 ± 0.0002 | 0.005 ± 0.001 | $2.88 \times 10^5$       | $2.88 \times 10^5$                    | 0.953             |

Table 17: MNIST ViT (50 tokens): pre-routed sparse attention (an 8-dimensional router selects  $k$  keys per query before attention) against post-hoc top- $k$  masking and full attention, 3 seeds. Attention FLOPs per block are given for the dense computation and with the sparsity exploited (post-hoc masking still computes the full score matrix and saves only the value aggregation; pre-routing computes scores for  $k$  keys only); the total ratio replaces the dense attention FLOPs of every block by the mode’s attention FLOPs and divides by the full model’s FLOPs.

| mode | keep ratio | k (keys/query) | seeds | test acc | ECE   | attn FLOPs/block (dense) | attn FLOPs/block (sparsity exploited) | total FLOPs ratio |
|------|------------|----------------|-------|----------|-------|--------------------------|---------------------------------------|-------------------|
| full | 1          | 65             | 1     | 0.7940   | 0.035 | $3.24 \times 10^6$       | $3.24 \times 10^6$                    | 1.000             |
| post | 0.25       | 17             | 1     | 0.8027   | 0.033 | $3.24 \times 10^6$       | $2.05 \times 10^6$                    | 0.980             |
| pre  | 0.25       | 17             | 1     | 0.8003   | 0.036 | $1.05 \times 10^6$       | $1.05 \times 10^6$                    | 0.964             |

Table 18: CIFAR-10 ViT (65 tokens): pre-routed sparse attention, one seed. Columns as in Table 17.

| group                                  | method                                                 | seeds | avg acc (final) | delta vs ref (pp) | forgetting      | retention     | ECE           | train FLOPs           |
|----------------------------------------|--------------------------------------------------------|-------|-----------------|-------------------|-----------------|---------------|---------------|-----------------------|
| sleep decay schedule                   | CLS2 + sparse fast (k-WTA 10%) (reference)             | 3     | 0.8703 ± 0.0059 | -                 | 0.1528 ± 0.0075 | 0.846 ± 0.008 | 0.080 ± 0.005 | $3.27 \times 10^{12}$ |
| sleep decay schedule                   | CLS2 + sleep decay: one-shot 10% at boundary           | 3     | 0.8735 ± 0.0015 | +0.32             | 0.1481 ± 0.0024 | 0.851 ± 0.002 | 0.077 ± 0.003 | $3.27 \times 10^{12}$ |
| sleep decay schedule                   | CLS2 + sleep decay: periodic gradual, total 10%        | 3     | 0.8729 ± 0.0020 | +0.27             | 0.1498 ± 0.0030 | 0.850 ± 0.003 | 0.072 ± 0.002 | $3.27 \times 10^{12}$ |
| sleep decay schedule                   | CLS2 + sleep decay: periodic gradual, total 3%         | 3     | 0.8750 ± 0.0087 | +0.47             | 0.1479 ± 0.0108 | 0.851 ± 0.011 | 0.074 ± 0.008 | $3.27 \times 10^{12}$ |
| sleep decay schedule                   | CLS2 + sleep decay: continuous weight decay 1e-4       | 3     | 0.8710 ± 0.0022 | +0.08             | 0.1523 ± 0.0043 | 0.847 ± 0.004 | 0.078 ± 0.001 | $3.27 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast width 256), no distillation (reference)     | 3     | 0.8547 ± 0.0032 | -                 | 0.1718 ± 0.0041 | 0.827 ± 0.004 | 0.093 ± 0.002 | $1.31 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + global logit distillation            | 3     | 0.7928 ± 0.0044 | -6.19             | 0.2509 ± 0.0049 | 0.748 ± 0.005 | 0.145 ± 0.005 | $1.31 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + local per-layer feature distillation | 3     | 0.8131 ± 0.0085 | -4.16             | 0.2209 ± 0.0108 | 0.778 ± 0.011 | 0.122 ± 0.012 | $1.31 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + local features + logits              | 3     | 0.7394 ± 0.0025 | -11.53            | 0.3062 ± 0.0066 | 0.691 ± 0.007 | 0.196 ± 0.013 | $1.31 \times 10^{12}$ |

Table 19: Split MNIST: sleep-decay schedules and distillation locality on the sparse redesign (CLS2,  $k$ -WTA 10% fast network), 3 seeds. Delta is the change in final average accuracy against the reference row of each group, in points.

| group                                  | method                                                 | seeds | avg acc (final) | delta vs ref (pp) | forgetting      | retention     | ECE           | train FLOPs           |
|----------------------------------------|--------------------------------------------------------|-------|-----------------|-------------------|-----------------|---------------|---------------|-----------------------|
| sleep decay schedule                   | CLS2 + sparse fast (k-WTA 10%) (reference)             | 3     | 0.8903 ± 0.0020 | -                 | 0.0662 ± 0.0024 | 0.931 ± 0.003 | 0.109 ± 0.022 | $1.18 \times 10^{13}$ |
| sleep decay schedule                   | CLS2 + sleep decay: one-shot 10% at boundary           | 3     | 0.8503 ± 0.0042 | -4.00             | 0.1125 ± 0.0042 | 0.882 ± 0.004 | 0.136 ± 0.004 | $1.18 \times 10^{13}$ |
| sleep decay schedule                   | CLS2 + sleep decay: periodic gradual, total 10%        | 3     | 0.8550 ± 0.0035 | -3.53             | 0.1075 ± 0.0028 | 0.888 ± 0.003 | 0.118 ± 0.008 | $1.18 \times 10^{13}$ |
| sleep decay schedule                   | CLS2 + sleep decay: periodic gradual, total 3%         | 3     | 0.8873 ± 0.0047 | -0.30             | 0.0718 ± 0.0047 | 0.925 ± 0.005 | 0.098 ± 0.011 | $1.18 \times 10^{13}$ |
| sleep decay schedule                   | CLS2 + sleep decay: continuous weight decay 1e-4       | 3     | 0.8850 ± 0.0058 | -0.53             | 0.0743 ± 0.0056 | 0.922 ± 0.006 | 0.094 ± 0.009 | $1.18 \times 10^{13}$ |
| distillation locality (fast width 256) | CLS2 (fast width 256), no distillation (reference)     | 3     | 0.8939 ± 0.0023 | -                 | 0.0636 ± 0.0019 | 0.934 ± 0.002 | 0.107 ± 0.010 | $4.75 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + global logit distillation            | 3     | 0.8782 ± 0.0064 | -1.58             | 0.0941 ± 0.0070 | 0.902 ± 0.007 | 0.120 ± 0.011 | $4.75 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + local per-layer feature distillation | 3     | 0.8516 ± 0.0047 | -4.23             | 0.1067 ± 0.0049 | 0.888 ± 0.005 | 0.178 ± 0.021 | $4.75 \times 10^{12}$ |
| distillation locality (fast width 256) | CLS2 (fast 256) + local features + logits              | 3     | 0.8355 ± 0.0079 | -5.85             | 0.1266 ± 0.0076 | 0.867 ± 0.008 | 0.172 ± 0.017 | $4.75 \times 10^{12}$ |

Table 20: Permuted MNIST: sleep-decay schedules and distillation locality, 3 seeds. Columns as in Table 19.

| density | arm                                                                     | seeds | active/sample | test acc            | ECE               | infer FLOPs        | cum. train FLOPs      | Jaccard same      | Jaccard diff      | coverage          |
|---------|-------------------------------------------------------------------------|-------|---------------|---------------------|-------------------|--------------------|-----------------------|-------------------|-------------------|-------------------|
| 0.2     | static: dense small (additive)                                          | 2     | 2,239,484     | 0.9093 $\pm$ 0.0000 | 0.035 $\pm$ 0.002 | 2.25 $\times 10^8$ | 6.74 $\times 10^{14}$ | -                 | -                 | -                 |
| 0.2     | static: prune-during, magnitude (global)                                | 2     | 2,232,870     | 0.9189 $\pm$ 0.0013 | 0.038 $\pm$ 0.000 | 3.45 $\times 10^8$ | 1.65 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.2     | dynamic, channel-level k-WTA input gating                               | 2     | 2,233,536     | 0.7339 $\pm$ 0.0072 | 0.021 $\pm$ 0.000 | 2.27 $\times 10^8$ | 1.36 $\times 10^{15}$ | 0.192 $\pm$ 0.010 | 0.127 $\pm$ 0.002 | 1.000 $\pm$ 0.000 |
| 0.2     | dynamic, connection-level, random selection                             | 2     | 2,238,912     | 0.7614 $\pm$ 0.0023 | 0.024 $\pm$ 0.003 | 2.25 $\times 10^8$ | 1.36 $\times 10^{15}$ | 0.111 $\pm$ 0.000 | 0.111 $\pm$ 0.000 | 1.000 $\pm$ 0.000 |
| 0.2     | dynamic, connection-level, local rule (abs(w) $\times$ input magnitude) | 2     | 2,238,912     | 0.9183 $\pm$ 0.0010 | 0.038 $\pm$ 0.000 | 2.25 $\times 10^8$ | 1.36 $\times 10^{15}$ | 0.613 $\pm$ 0.005 | 0.462 $\pm$ 0.008 | 0.429 $\pm$ 0.003 |
| 0.05    | static: dense small (additive)                                          | 3     | 562,229       | 0.8866 $\pm$ 0.0018 | 0.030 $\pm$ 0.002 | 5.58 $\times 10^7$ | 1.67 $\times 10^{14}$ | -                 | -                 | -                 |
| 0.05    | static: train-then-prune + finetune                                     | 3     | 558,218       | 0.9227 $\pm$ 0.0006 | 0.035 $\pm$ 0.001 | 2.71 $\times 10^8$ | 3.74 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.05    | static: RigL                                                            | 3     | 558,213       | 0.9097 $\pm$ 0.0016 | 0.034 $\pm$ 0.001 | 1.39 $\times 10^8$ | 4.17 $\times 10^{14}$ | -                 | -                 | -                 |
| 0.05    | static: prune-during, magnitude (ERK)                                   | 3     | 558,213       | 0.9203 $\pm$ 0.0009 | 0.034 $\pm$ 0.001 | 1.39 $\times 10^8$ | 1.43 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.05    | static: prune-during, magnitude (global)                                | 3     | 558,218       | 0.9196 $\pm$ 0.0023 | 0.034 $\pm$ 0.001 | 1.85 $\times 10^8$ | 1.33 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.05    | dynamic, channel-level k-WTA input gating                               | 2     | 569,536       | 0.4486 $\pm$ 0.0036 | 0.016 $\pm$ 0.003 | 5.74 $\times 10^7$ | 9.83 $\times 10^{14}$ | 0.083 $\pm$ 0.002 | 0.052 $\pm$ 0.001 | 0.916 $\pm$ 0.008 |
| 0.05    | dynamic, connection-level, random selection                             | 2     | 564,160       | 0.5229 $\pm$ 0.0005 | 0.029 $\pm$ 0.001 | 5.91 $\times 10^7$ | 9.86 $\times 10^{14}$ | 0.026 $\pm$ 0.000 | 0.026 $\pm$ 0.000 | 1.000 $\pm$ 0.000 |
| 0.05    | dynamic, connection-level, local rule (abs(w) $\times$ input magnitude) | 2     | 564,160       | 0.8983 $\pm$ 0.0011 | 0.039 $\pm$ 0.003 | 5.91 $\times 10^7$ | 9.86 $\times 10^{14}$ | 0.439 $\pm$ 0.003 | 0.274 $\pm$ 0.002 | 0.207 $\pm$ 0.003 |
| 0.005   | static: dense small (additive)                                          | 3     | 55,872        | 0.8131 $\pm$ 0.0026 | 0.021 $\pm$ 0.001 | 6.12 $\times 10^6$ | 1.84 $\times 10^{13}$ | -                 | -                 | -                 |
| 0.005   | static: train-then-prune + finetune                                     | 3     | 55,822        | 0.8364 $\pm$ 0.0015 | 0.016 $\pm$ 0.001 | 5.27 $\times 10^7$ | 3.41 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.005   | static: RigL                                                            | 3     | 55,820        | 0.8486 $\pm$ 0.0030 | 0.026 $\pm$ 0.001 | 1.38 $\times 10^7$ | 4.15 $\times 10^{13}$ | -                 | -                 | -                 |
| 0.005   | static: prune-during, magnitude (ERK)                                   | 3     | 55,820        | 0.8783 $\pm$ 0.0022 | 0.024 $\pm$ 0.001 | 1.38 $\times 10^7$ | 1.18 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.005   | static: prune-during, magnitude (global)                                | 3     | 55,822        | 0.8896 $\pm$ 0.0003 | 0.027 $\pm$ 0.001 | 3.40 $\times 10^7$ | 1.09 $\times 10^{15}$ | -                 | -                 | -                 |
| 0.005   | dynamic, channel-level k-WTA input gating                               | 2     | 69,952        | 0.3404 $\pm$ 0.0296 | 0.021 $\pm$ 0.001 | 1.34 $\times 10^7$ | 8.80 $\times 10^{14}$ | 0.058 $\pm$ 0.007 | 0.036 $\pm$ 0.002 | 0.285 $\pm$ 0.021 |
| 0.005   | dynamic, connection-level, random selection                             | 2     | 63,424        | 0.1824 $\pm$ 0.0026 | 0.040 $\pm$ 0.002 | 9.32 $\times 10^6$ | 8.75 $\times 10^{14}$ | 0.002 $\pm$ 0.000 | 0.003 $\pm$ 0.000 | 0.633 $\pm$ 0.000 |
| 0.005   | dynamic, connection-level, local rule (abs(w) $\times$ input magnitude) | 2     | 63,424        | 0.7768 $\pm$ 0.0015 | 0.019 $\pm$ 0.001 | 9.32 $\times 10^6$ | 8.75 $\times 10^{14}$ | 0.333 $\pm$ 0.009 | 0.246 $\pm$ 0.006 | 0.033 $\pm$ 0.001 |

Table 21: Input-dependent (per-sample) pruning of a CIFAR-10 ResNet-18 against fixed masks at the same per-sample budget. Active weights are the per-sample expectation; Jaccard values are overlaps of the last-stage masks of two test images of the same or of different classes (200 images; a random mask gives 0.026); union coverage is the fraction of that stage’s connections used by at least one of the 200 images.

| $n_e$ | seeds | STDP acc            | MLP acc             | SNN SOP/img        | MLP FLOPs/img      | SNN train ops         | MLP train ops         | FLOPs                 | SNN to 80%      | MLP to 80%   | $w < 1\% w_{\max}$ | time (s)    |
|-------|-------|---------------------|---------------------|--------------------|--------------------|-----------------------|-----------------------|-----------------------|-----------------|--------------|--------------------|-------------|
| 100   | 1     | 0.3697              | 0.9474              | $1.08 \times 10^5$ | $1.59 \times 10^5$ | $1.15 \times 10^{10}$ | $2.86 \times 10^{10}$ | $2.86 \times 10^{10}$ | -               | 4992         | 0.549              | 550         |
| 400   | 3     | $0.8108 \pm 0.0028$ | $0.9649 \pm 0.0021$ | $4.33 \times 10^5$ | $6.35 \times 10^5$ | $4.60 \times 10^{10}$ | $1.14 \times 10^{11}$ | $1.14 \times 10^{11}$ | $38189 \pm 492$ | $4992 \pm 0$ | $0.577 \pm 0.000$  | $561 \pm 9$ |
| 1600  | 1     | 0.7289              | 0.9682              | $1.74 \times 10^6$ | $2.54 \times 10^6$ | $1.83 \times 10^{11}$ | $4.57 \times 10^{11}$ | $4.57 \times 10^{11}$ | -               | 4992         | 0.287              | 576         |

Table 22: Unsupervised STDP (Diehl & Cook) against a backprop MLP of the same width trained on the same samples. SOP = synaptic operations (event-driven, accumulated). No energy comparison is made: everything runs on a GPU.